

CS 59300 – Algorithms for Data Science

Classical and Quantum approaches

Lecture 4 (09/09)

Tensor Methods (IV)

https://ruizhezhang.com/course_fall_2025.html

Recap

We've seen several tensor decomposition algorithms:

- Jennrich's algorithm (simultaneous diagonalization)
- Tensor power method
- Alternating least squares
- Flattening-based higher-order tensor decomposition

However, the equations for tensors are too long and involves too many indices and \sum 's, e.g.,

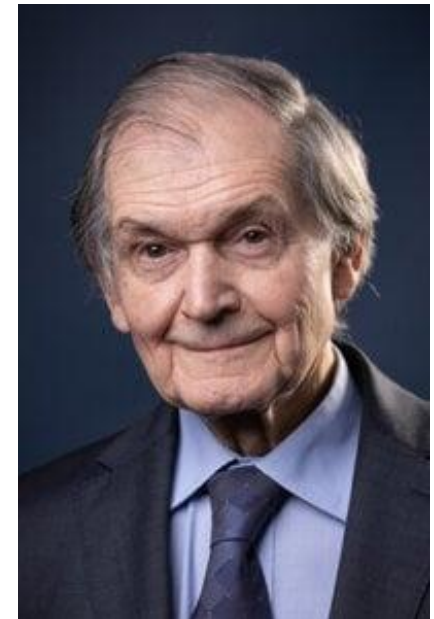
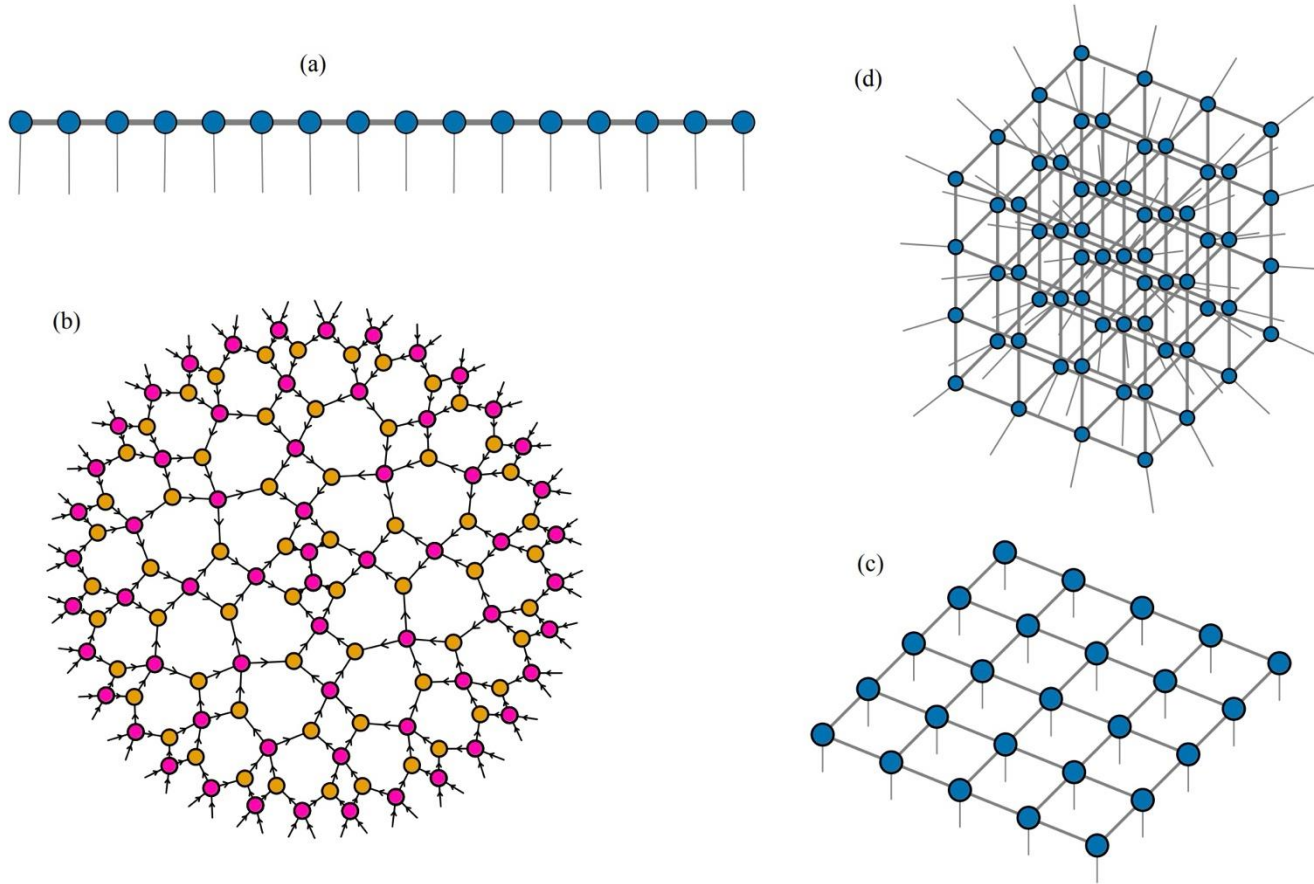
$$T'_{abc} = \sum_{a'b'c'} \sum_i \lambda_i (u_i)_{a'} (u_i)_{b'} (u_i)_{c'} W_{a'a} W_{b'b} W_{c'c}$$

Today: we'll see a **diagrammatic language** for tensors---tensor networks

Today's plan

- Tensor diagram notations
- Tensor networks
- Quantum application: classical simulation of quantum circuits
- Tensor network algorithms

Tensor diagrams



Roger Penrose

Tensor diagram rule #1

Tensors are notated by shapes (usually filled or shaded), and tensor indices are notated by lines (or “legs”) emanating from these shapes.

vector

v_j



matrix

M_{ij}



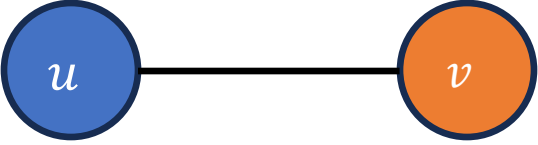
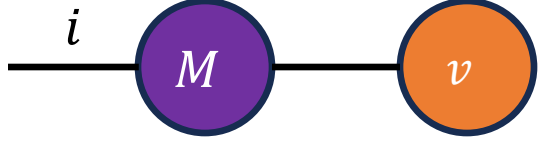
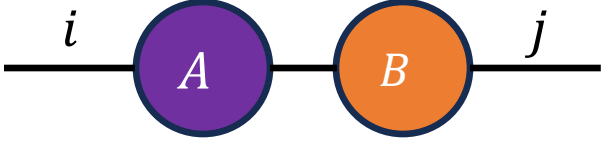
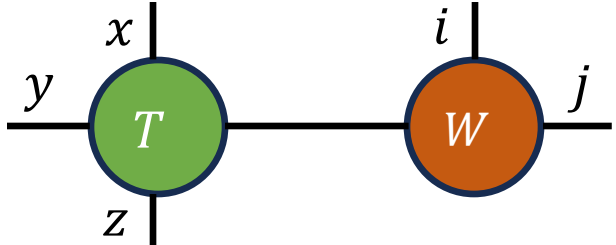
3-index
tensor

T_{ijk}



Tensor diagram rule #2

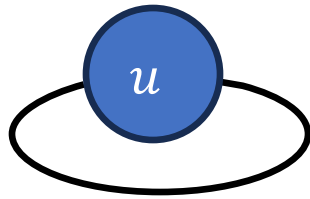
Connecting two index lines implies a **contraction**, or summation over the connected indices.

	=	$\sum_i u_i v_i = \langle u, v \rangle$
	=	$\sum_i M_{ji} v_i = Mv$
	=	$\sum_k A_{ik} B_{kj} = AB$
	=	$\sum_\alpha T_{xyz\alpha} W_{ij\alpha}$

Tensor diagram rule #2

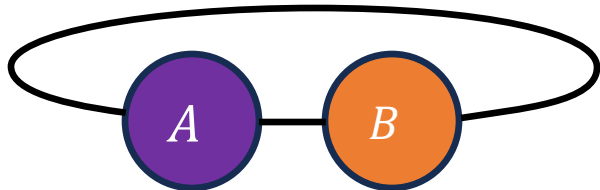
Connecting two index lines implies a **contraction**, or summation over the connected indices.

(can also involve loops)



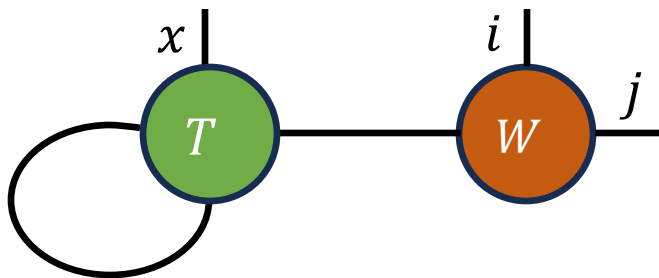
=

$$\sum_i u_i u_i = \|u\|^2$$



=

$$\sum_{ij} A_{ij} B_{ji} = \text{tr}[AB]$$



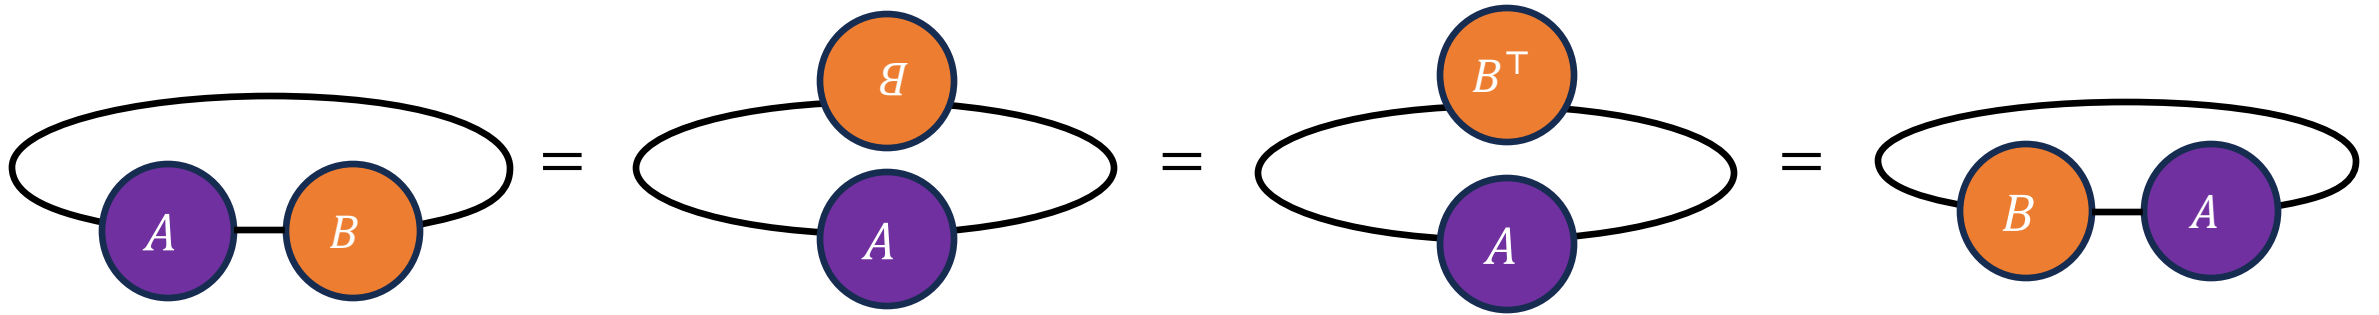
=

$$\sum_{\alpha, \beta} T_{x\beta\beta\alpha} W_{ij\alpha}$$

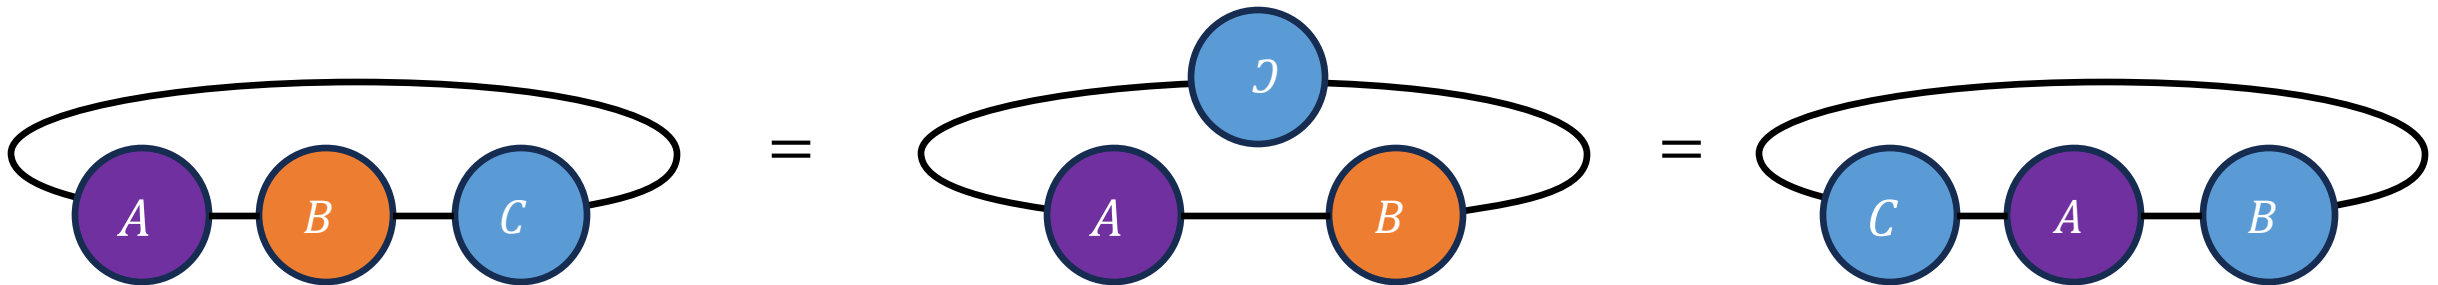
partial trace

Diagrammatic proof of the trace identity

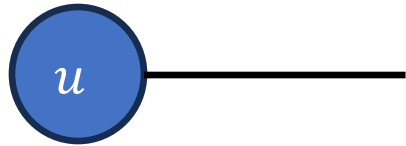
$$\text{tr}[AB] = \text{tr}[BA]$$



$$\text{tr}[ABC] = \text{tr}[CAB]$$

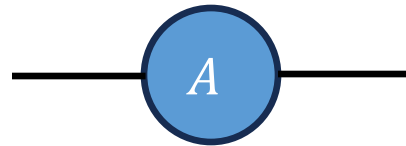


Tensor products



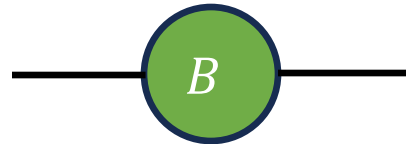
=

$$u \otimes v$$

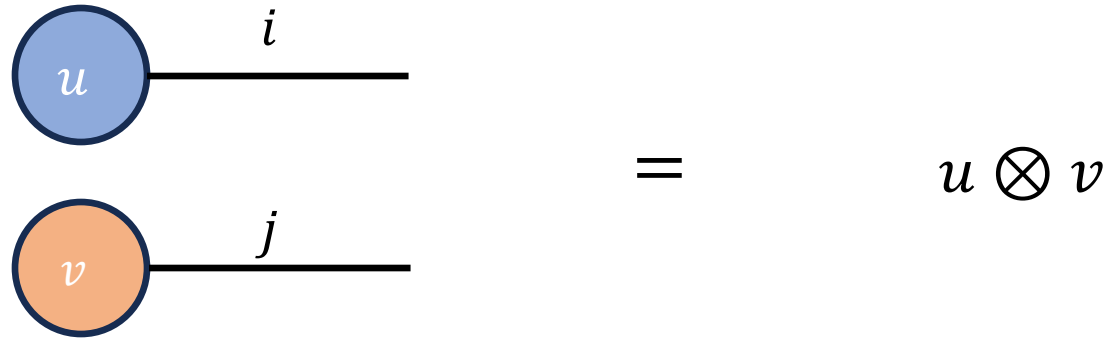


=

$$A \otimes B$$



Grouping indices



The diagram illustrates the grouping of indices. On the left, there are two separate components: a blue circle labeled u with a horizontal line extending to the right labeled i , and an orange circle labeled v with a horizontal line extending to the right labeled j . These are followed by an equals sign, and then the tensor product expression $u \otimes v$.

$$\begin{array}{c} \text{blue circle } u \text{ --- } i \\ \text{orange circle } v \text{ --- } j \end{array} = u \otimes v$$

Grouping indices

Diagrammatic equation showing the contraction of a vector v (represented by a brown circle) with an index j (represented by a line) to form a new index ij (represented by a green circle). The result is the vector $\text{vec}(u \otimes v)$.

$$\text{brown circle } v \text{ --- } j = \text{green circle } ij = \text{vec}(u \otimes v)$$

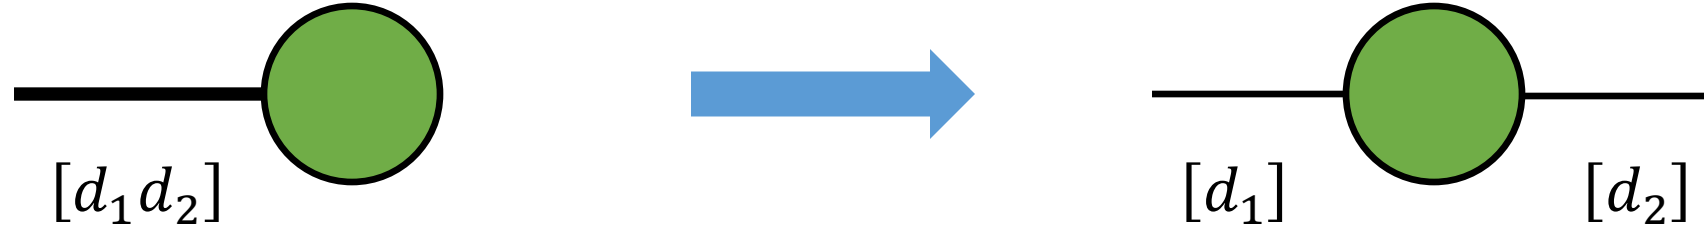
Diagrammatic equation showing the contraction of a tensor A (represented by a purple circle) with indices i and j (represented by lines) to form a new index ij (represented by a purple circle). The result is the vector $\text{vec}(A)$.

$$\text{purple circle } A \text{ with lines } i \text{ and } j = \text{purple circle } ij = \text{vec}(A)$$

Diagrammatic equation showing the contraction of two tensors (represented by blue circles) with indices i and j (represented by lines) to form a new tensor (represented by a blue circle). The result is the tensor $\text{vec}(A)$.

$$\text{blue circle } A \text{ with lines } i \text{ and } j = \text{blue circle } ij = \text{vec}(A)$$

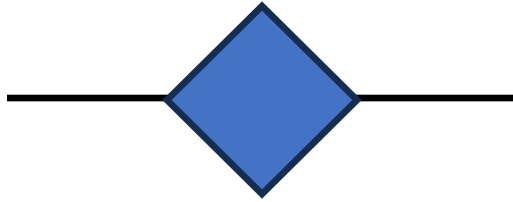
Splitting indices



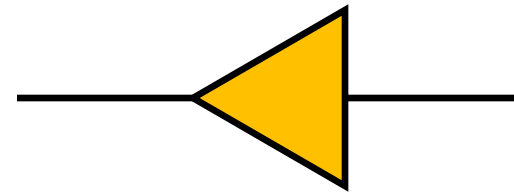
Some special notations



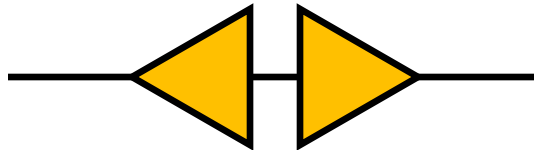
identity



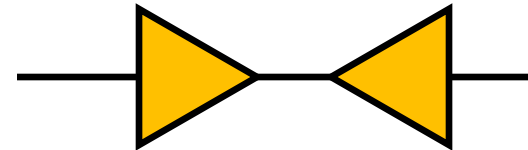
diagonal matrix



isometry



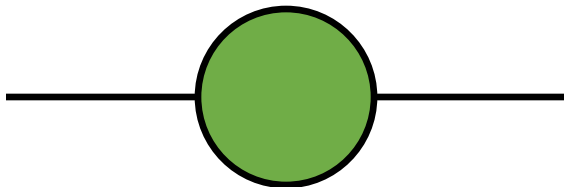
=



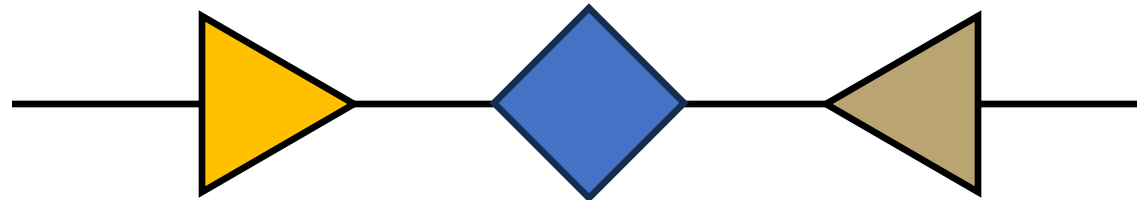
≠



SVD $A = U\Sigma V^\dagger$

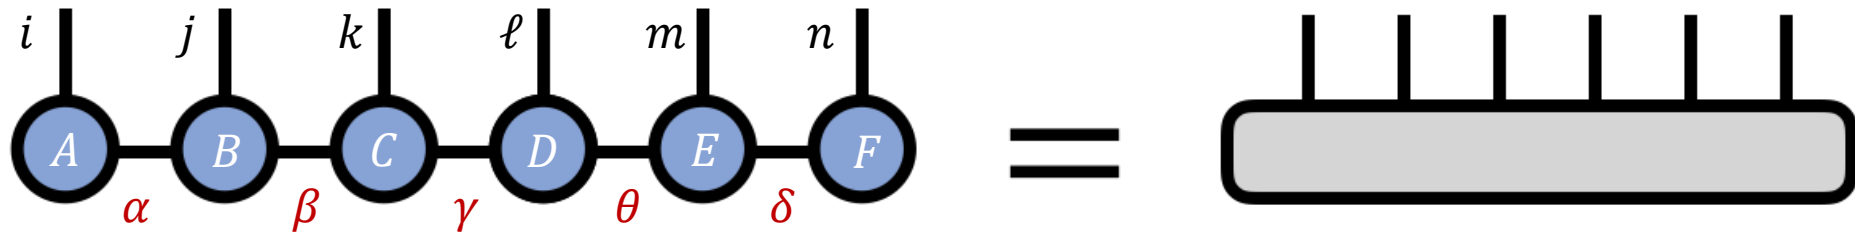


=



Tensor networks

Tensor networks are **factorizations** of very large tensors into networks of smaller tensors

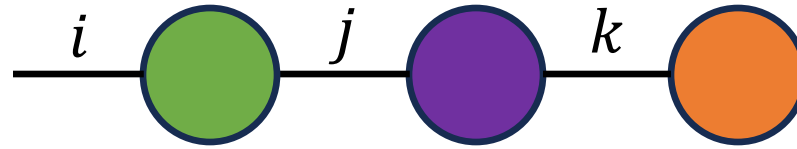


$$\sum_{\alpha\beta\gamma\theta\delta} A_{i\alpha} B_{\alpha j\beta} C_{\beta k\gamma} D_{\gamma\ell\theta} E_{\theta m\delta} F_{\delta n} = T_{ijklmn}$$

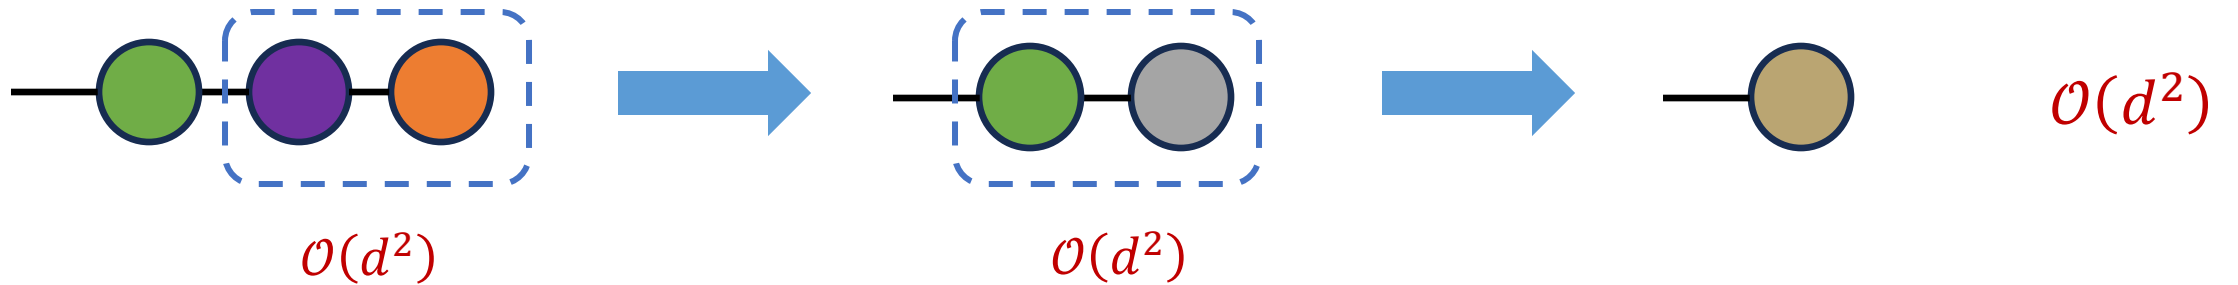
Q: Does the order of contractions matter?

A: *No, mathematically. Yes, computationally!*

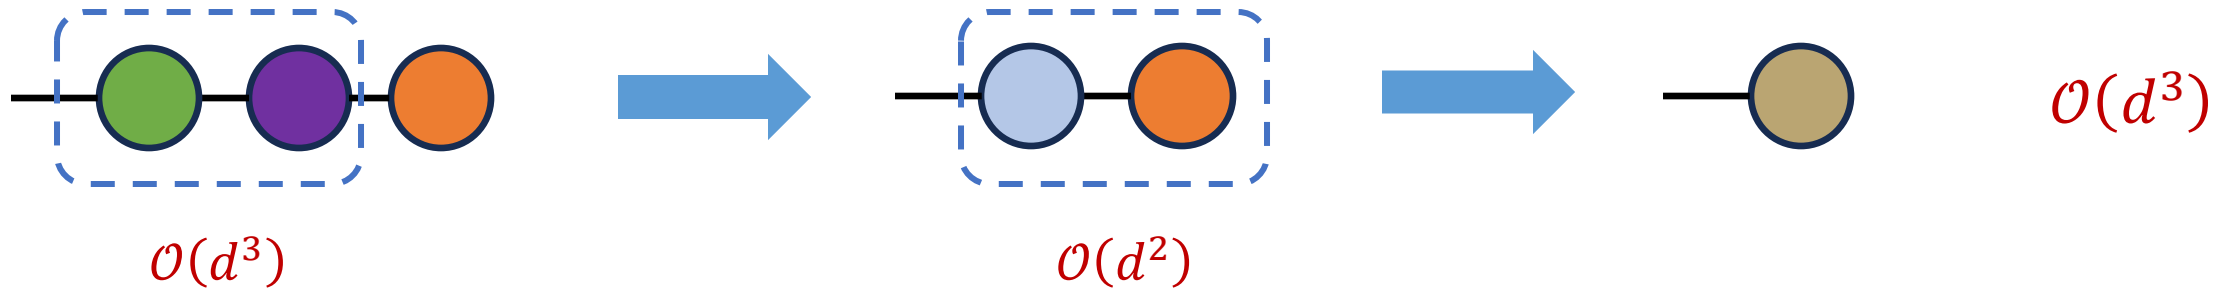
Computational costs of tensor contraction



- Contraction path 1:



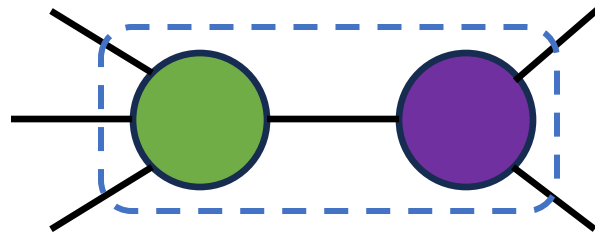
- Contraction path 2:



Computational costs of tensor contraction

Fact. The computational cost for contracting an edge with dimension d is:

$$\mathcal{O}\left(d \cdot \prod_{k:\text{open edges}} d_k\right)$$

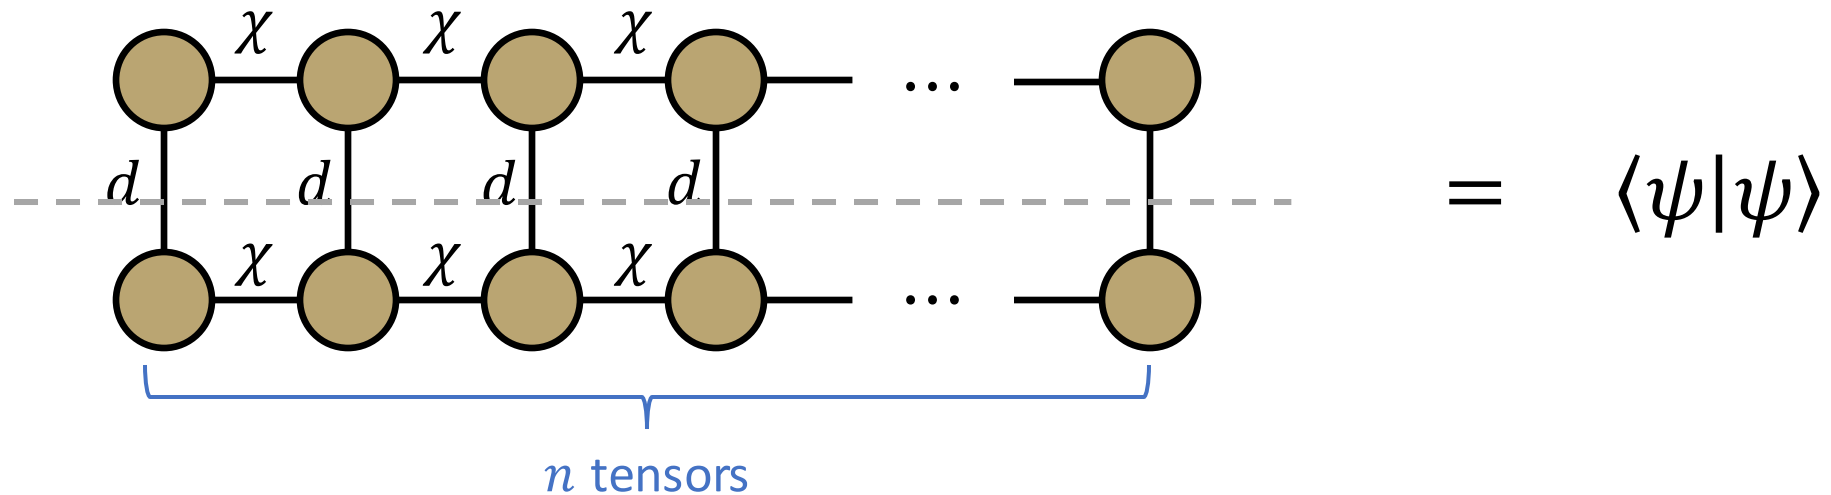


$$\mathcal{O}(d^6)$$

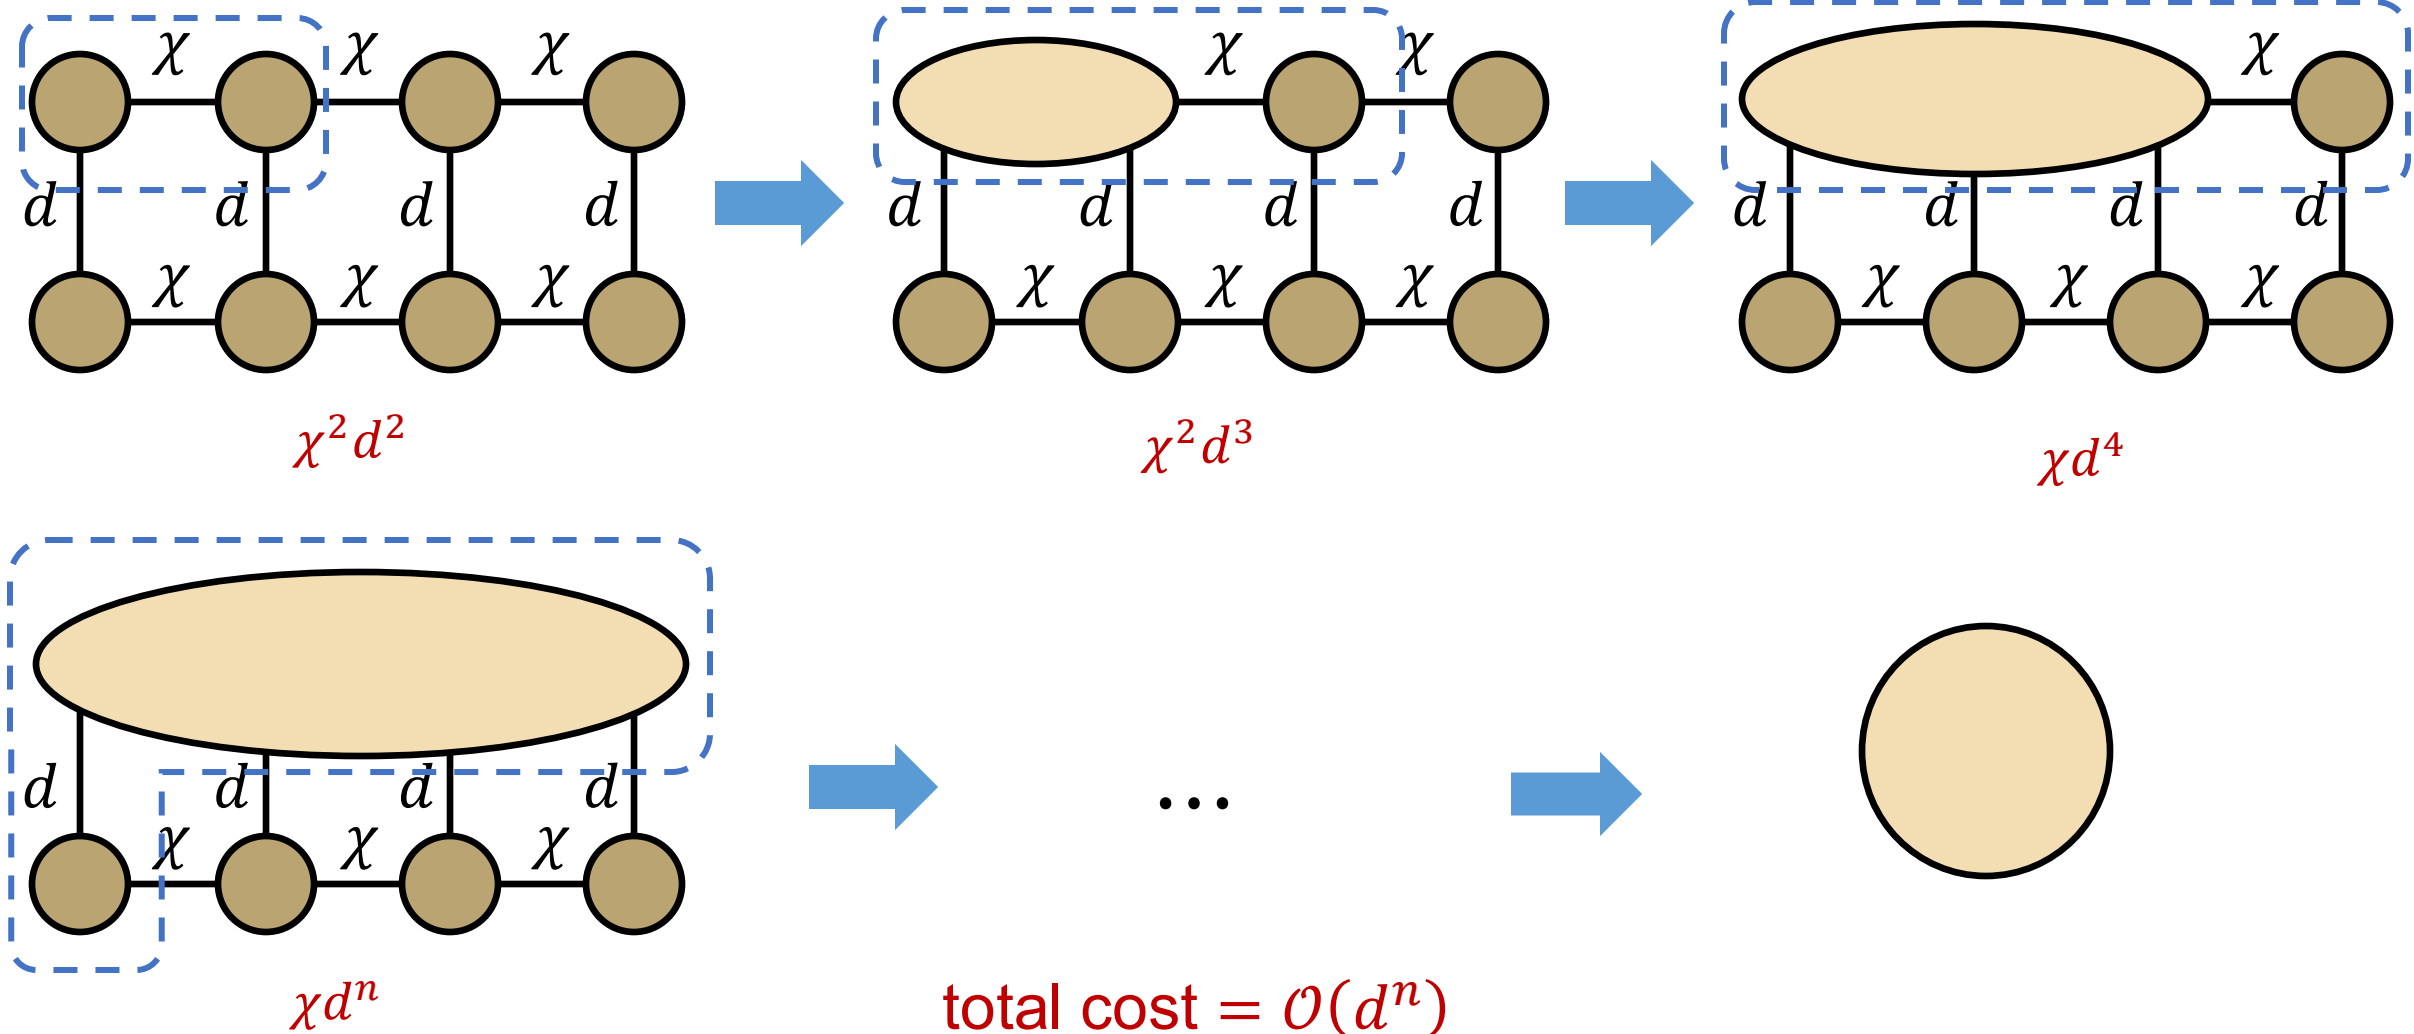
Computational costs of tensor contraction

Fact. The computational cost for contracting an edge with dimension d is:

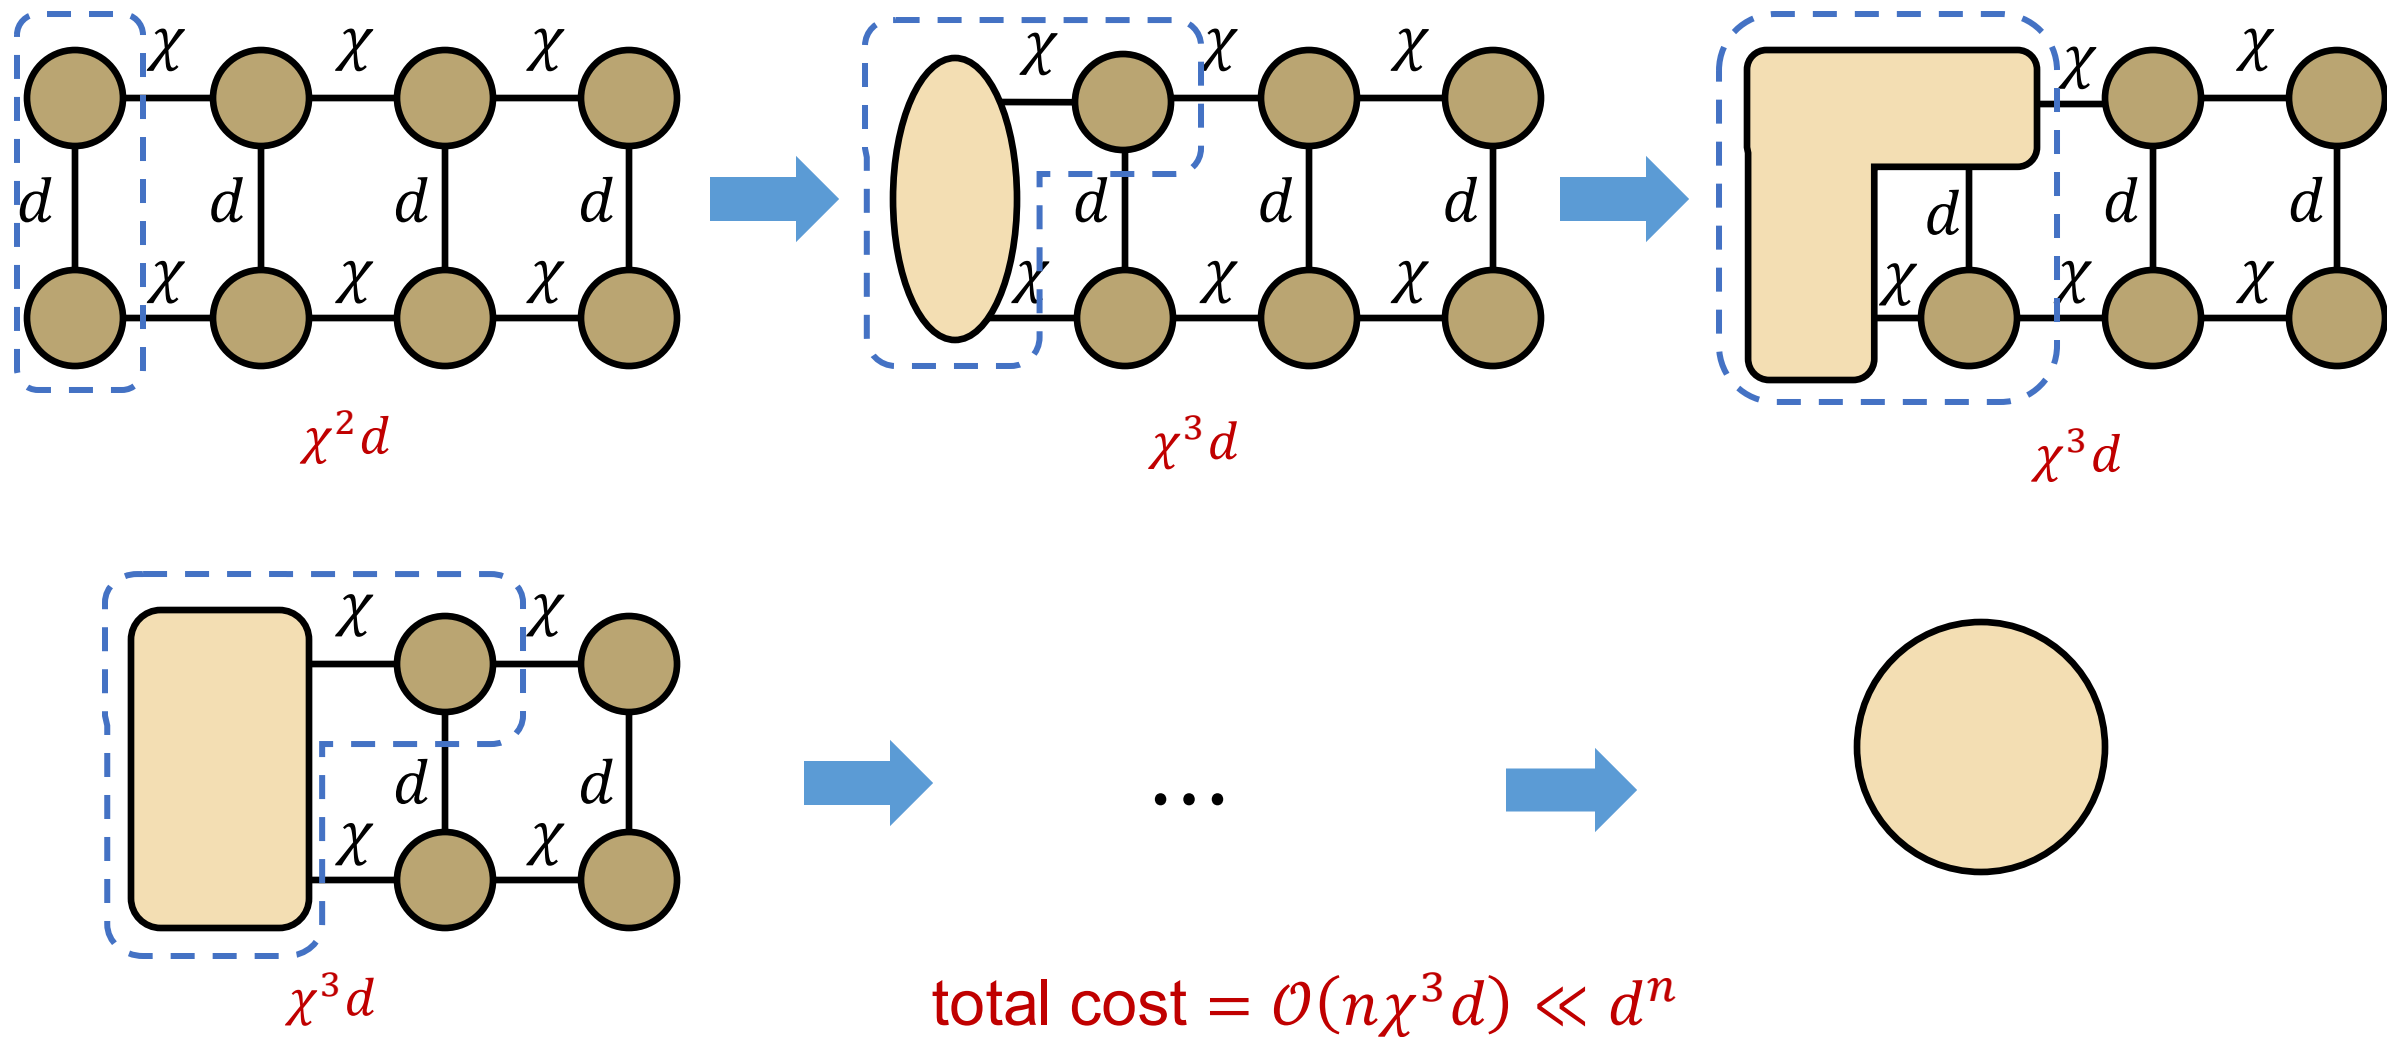
$$\mathcal{O}\left(d \cdot \prod_{k:\text{open edges}} d_k\right)$$



Contracting a ladder: left-to-right then top-to-bottom



Contracting a ladder: staggered



Tensor network contraction is hard

Theorem. Tensor network contraction is **#P**-complete.

Proof.

- Encode **#SAT** as a tensor network: its contraction equals the number of satisfying assignments

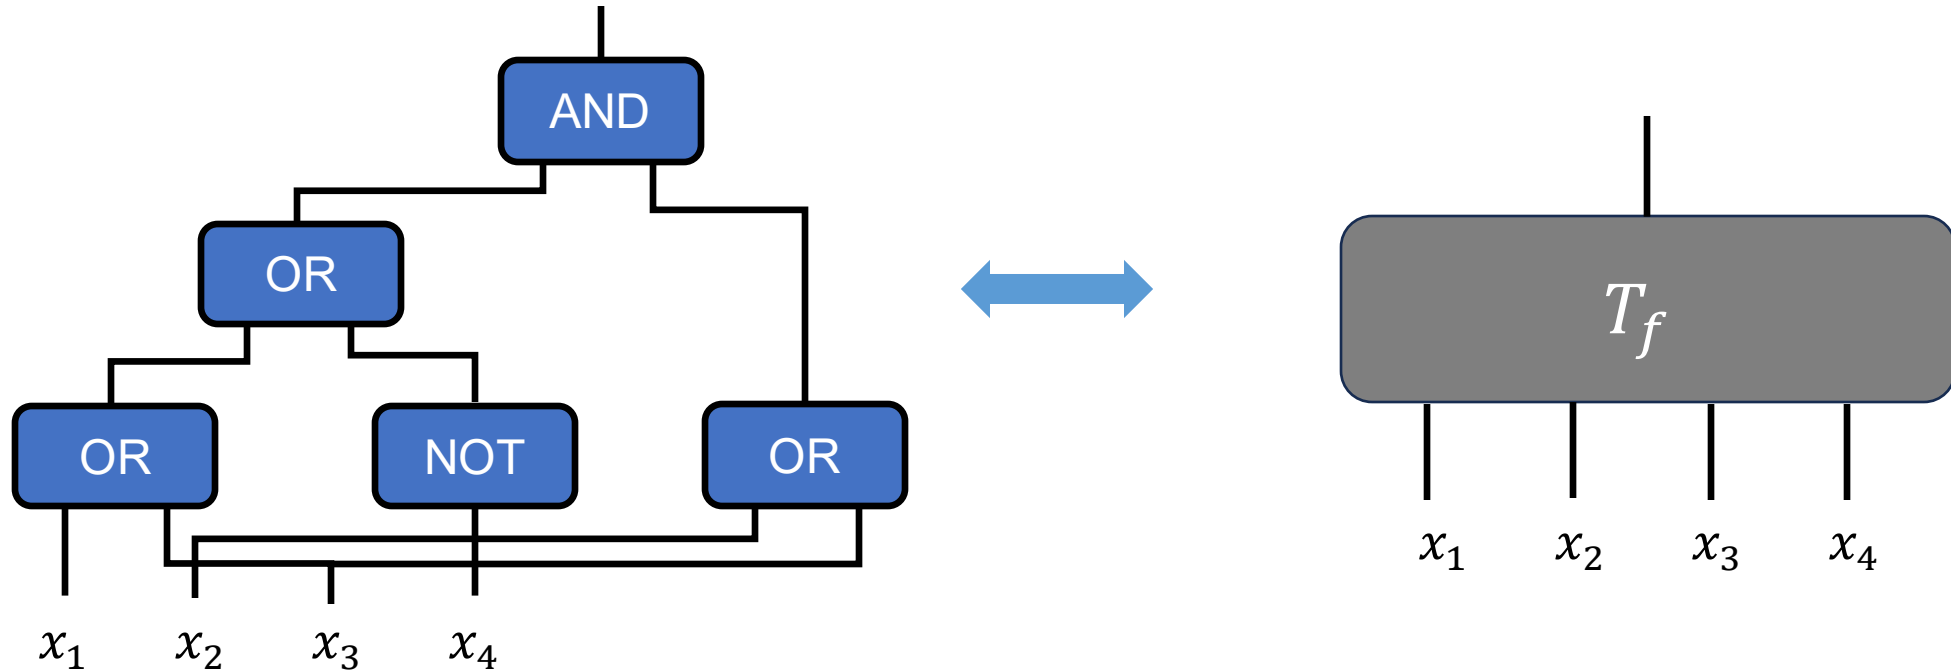
$$\forall f \in \{0,1\}^n \rightarrow \{0,1\} \quad \longrightarrow \quad T_f := \sum_{x \in \{0,1\}^n} |f(x)\rangle \langle x| \quad \in \quad \{0,1\}^{\otimes n} \otimes \{0,1\}$$



Tensor network contraction is hard

Theorem. Tensor network contraction is **#P**-complete.

Example: $f(x) = (x_1 \vee x_3 \vee \overline{x_4}) \wedge (x_2 \vee \overline{x_3})$



Tensor network contraction is hard

Theorem. Tensor network contraction is **#P**-complete.

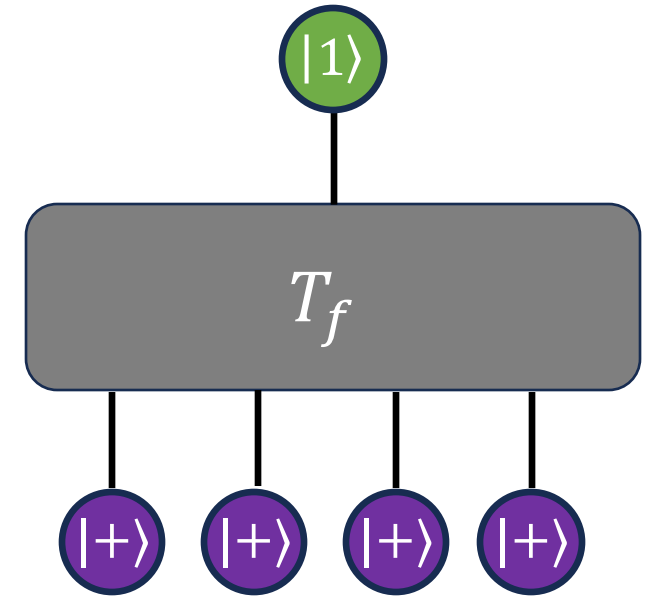
- How to count the solutions?

- Let $|+\rangle := |0\rangle + |1\rangle$ $|+\rangle = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$

- $|+\rangle^{\otimes n} = (|0\rangle + |1\rangle)^{\otimes n} = \sum_{x \in \{0,1\}^n} |x\rangle$

$$|+\rangle \otimes |+\rangle = |00\rangle + |01\rangle + |10\rangle + |11\rangle$$

- $T_f(|+\rangle, \dots, |+\rangle, |1\rangle) = \#\text{satisfiable assignments for } f$
- If there exists a polynomial time algorithm for tensor contraction, then there is also a polynomial time algorithm for #SAT, which is **#P**-complete



Contraction complexity

A tensor network can be described as an undirected graph $G = (V, E)$:

- Vertices \longleftrightarrow tensors, edges \longleftrightarrow indices
- Contraction of an edge e removes e and replaces its end vertices (or vertex) with a single vertex

A **contraction ordering** π is an ordering of all edges:

$$\{\pi_1, \pi_2, \dots, \pi_{|E|}\} = E$$

- The complexity of π is the **maximum degree of a merged vertex** during the contraction process

Contraction complexity of G , denoted by $\text{cc}(G)$, is the minimum complexity of a contraction ordering

The cost of contracting the TN is $\sim d^{\mathcal{O}(\text{cc}(G))}$ or $\exp\left(\mathcal{O}(\text{cc}(G))\right)$ (for constant dimensions)

How to determine $\text{cc}(G)$?

Contraction complexity: basic properties

Claim. It holds that

$$\Delta(G) - 1 \leq cc(G) \leq |E| - 1$$

where $\Delta(G)$ is the maximum degree of a vertex in G

Proof.

- Since a merged vertex cannot be connected to more than $|E| - 1$ edges, so $cc(G) \leq |E| - 1$
- When any edge incident to a vertex of degree $\Delta(G)$ is removed, the resulting merged vertex is incident to at least $\Delta(G) - 1$ edges. Thus, $\Delta(G) - 1 \leq cc(G)$

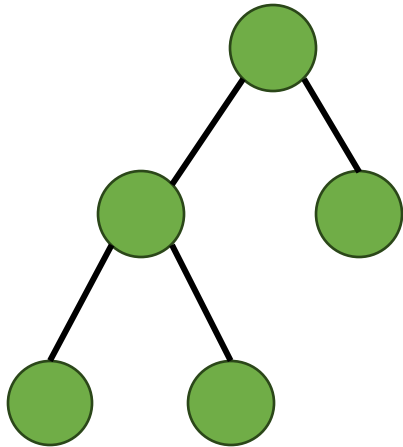
Contraction complexity

Theorem.

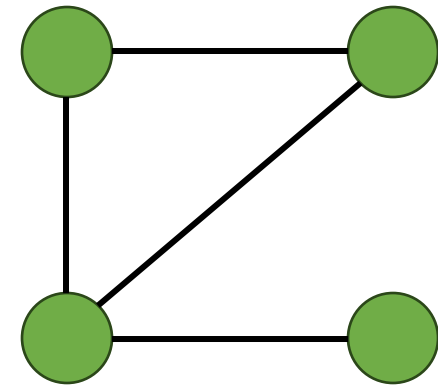
Given a tensor network with N tensors and underlying graph G , the contraction time is

$$\mathcal{O}(N \exp(\text{tw}(G)))$$

where $\text{tw}(G)$ is the **tree-width** of G (assuming all indices are of $\mathcal{O}(1)$ dimensions)

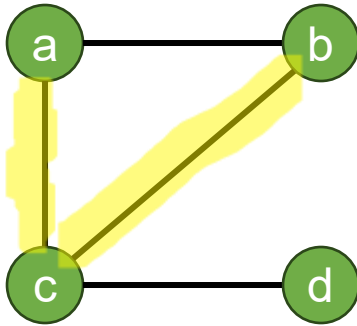


Tree decomposition is a way to measure how **tree-like** a graph is

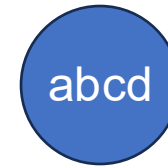


Tree decomposition

How to transform this graph to a tree?

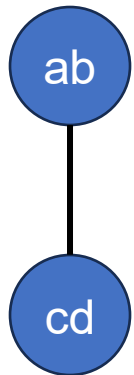


- Trivial:



Group of vertices
(bag)

- Slightly non-trivial:



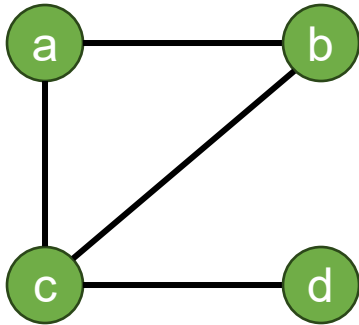
Does not reflect the structure of
the original graph

We require:

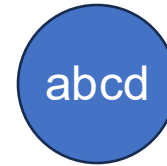
$\forall e = (i, j) \in E$, there exists a bag in
the tree that contains both i and j

Tree decomposition

How to transform this graph to a tree?

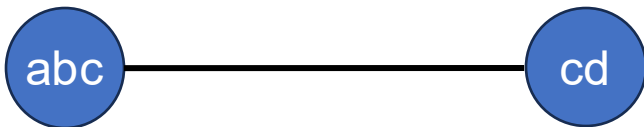


- Trivial:



$$\max_{u \in T} |L_u| = 4$$

- Another choice:



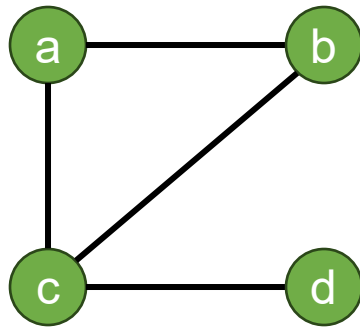
$$\max_{u \in T} |L_u| = 3$$

Check-list

- Tree? ☒
- Contains all vertices? ☒
- Contains all edges? ☒
- Minimize the max bag size

Tree decomposition

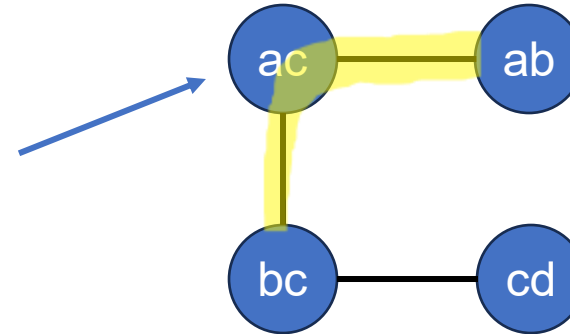
How to transform this graph to a tree?



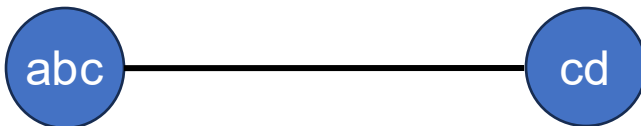
forget the
vertex b



- Smaller bag?



- Another choice:



$$\max_{u \in T} |L_u| = 3$$

Check-list

- Tree? ☒
- Contains all vertices? ☒
- Contains all edges? ☒
- Minimize the max bag size

Tree decomposition

A tree decomposition of a graph $G = (V, E)$ is a tree of N nodes u_1, \dots, u_N , with a set $L_{u_i} \subset V$ corresponding to each node u_i , such that:

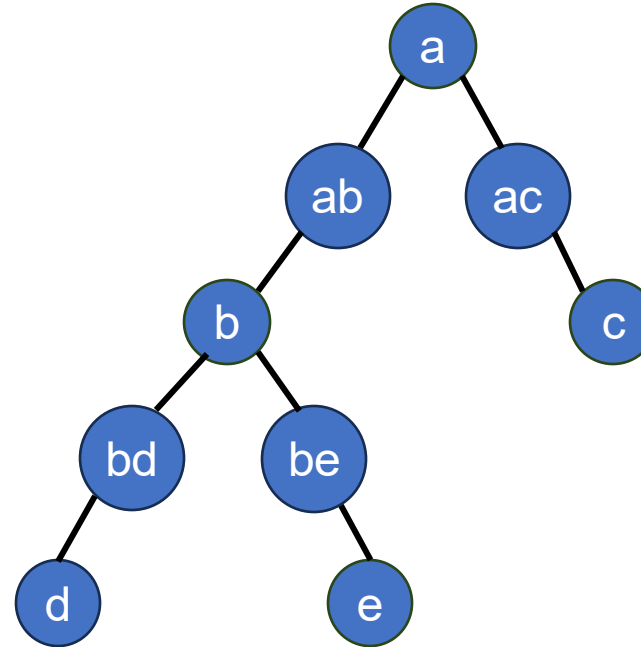
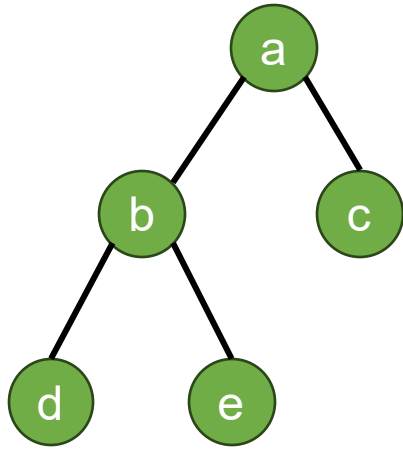
1. **Vertex coverage:** every vertex of G belongs to at least one set
2. **Edge coverage:** for every edge in G , there is a set containing both its endpoints
3. **Consistency:** for every vertex v in G , the set of bags containing v induces a connected subgraph
Formally, if $v \in L_{u_i} \cap L_{u_j}$, then $v \in L_{u_k}$ for all nodes u_k on the $u_i \rightarrow u_j$ path

The **width** of a tree decomposition is the **maximum size of a bag**

The **tree-width** of a graph G is the **minimum width** of any tree decomposition of G **minus one**

Tree decomposition

Tree has tree-width **one**:



Contraction complexity

Theorem (Markov-Shi, 2008).

Given a tensor network with N tensors and underlying graph G , the contraction time is

$$\mathcal{O}(N \exp(\text{tw}(G)))$$

where $\text{tw}(G)$ is the **tree-width** of G

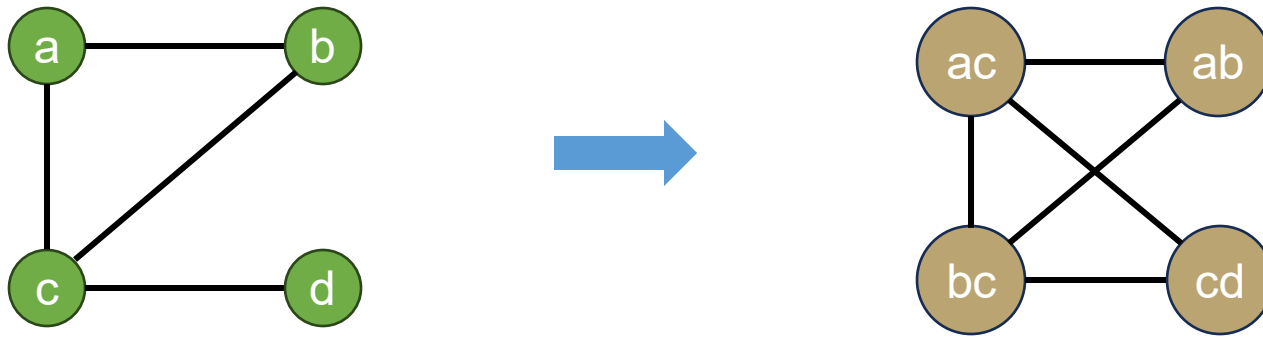
Proof strategy:

1. Prove that $\text{cc}(G) = \text{tw}(G^*)$, where G^* is the **line graph** of G
2. Use $\text{tw}(G)$ to bound $\text{tw}(G^*)$

Line graph

The **line graph** of $G = (V, E)$, denoted as G^* , has vertex set $V(G^*) = E$, and edge set

$$E(G^*) = \{(e_1, e_2) : e_1, e_2 \in E \text{ and } e_1 \cap e_2 \neq \emptyset\}$$



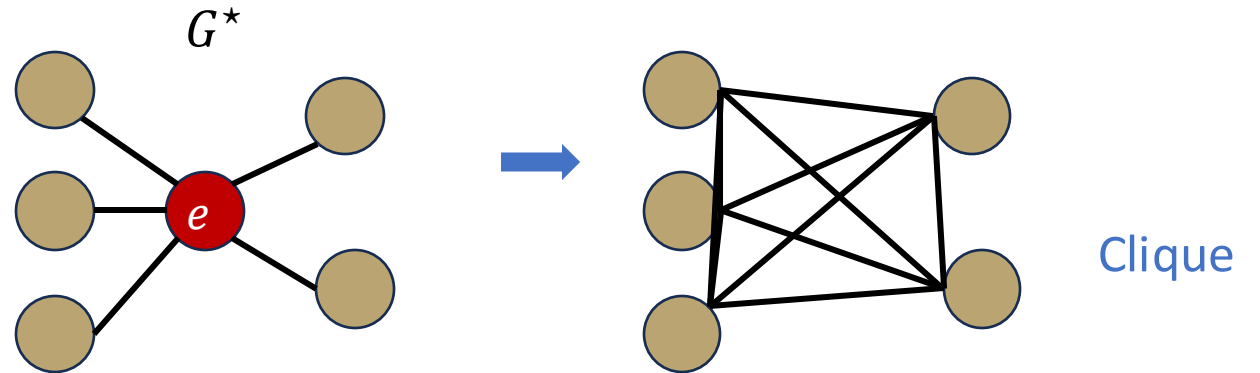
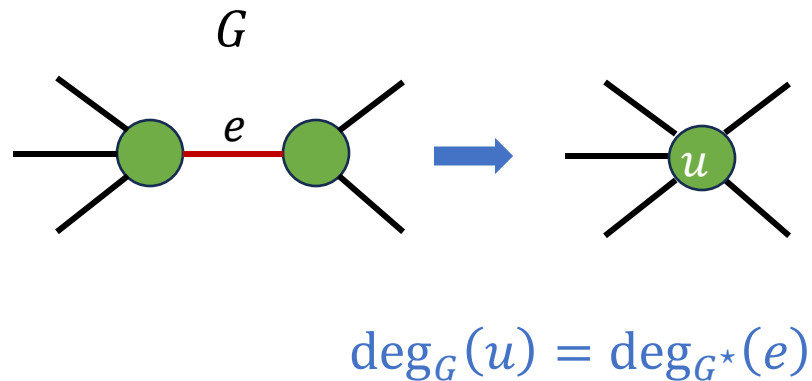
Contraction complexity equals tree-width

Proposition.

For any graph $G = (V, E)$, $cc(G) = tw(G^*)$. Furthermore, given a tree decomposition of G^* of width d , there is a deterministic algorithm that outputs a contraction ordering π with $cc(\pi) \leq d$ in polynomial time.

Proof ideas:

- When we contract an edge $e \in E$, it corresponds to removing a vertex in G^*



Contraction complexity equals tree-width

Proposition.

For any graph $G = (V, E)$, $cc(G) = tw(G^*)$. Furthermore, given a tree decomposition of G^* of width d , there is a deterministic algorithm that outputs a contraction ordering π with $cc(\pi) \leq d$ in polynomial time.

Proof ideas:

- When we contract an edge $e \in E$, it corresponds to removing a vertex in G^* , and connect its neighbors as a clique
- The degree of the merged vertex in G = the degree of e in G^*
- These operations exactly correspond to the **elimination width (or induced width)** of G^*
- Graph theory tells that **elimination width = tree-width**

Tree-width of G and G^\star

Lemma. For any graph G of maximum degree $\Delta(G)$,

$$(\text{tw}(G) - 1)/2 \leq \text{tw}(G^\star) \leq \Delta(G)(\text{tw}(G) + 1) + 1$$

- So, for a bounded-degree graph G , $\text{tw}(G) \simeq \text{tw}(G^\star)$

How to find the tree decomposition?

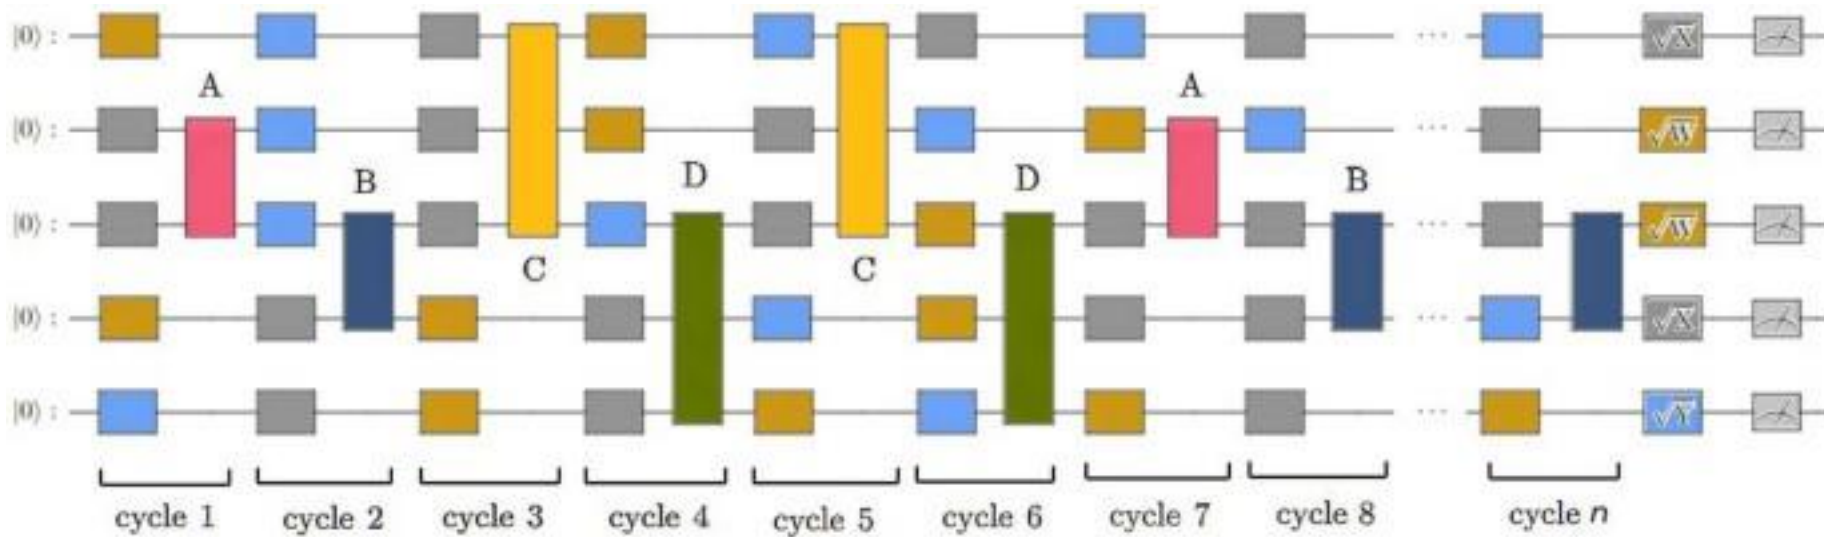
- **Robertson-Seymour:** There is a deterministic algorithm that given a graph G outputs a tree decomposition of G of width $\mathcal{O}(\text{tw}(G))$ in time $|V|^{\mathcal{O}(1)} \exp(\mathcal{O}(\text{tw}(G)))$.

Tensor network contraction algorithm

1. Apply the Robertson-Seymour algorithm to compute a tree decomposition
2. Find the contraction ordering π using the proposition
3. Contract the tensor network according to π

The first and the third steps take $N^{\mathcal{O}(1)} \exp\left(\mathcal{O}(\text{tw}(G))\right)$ time, and the second step is cheap

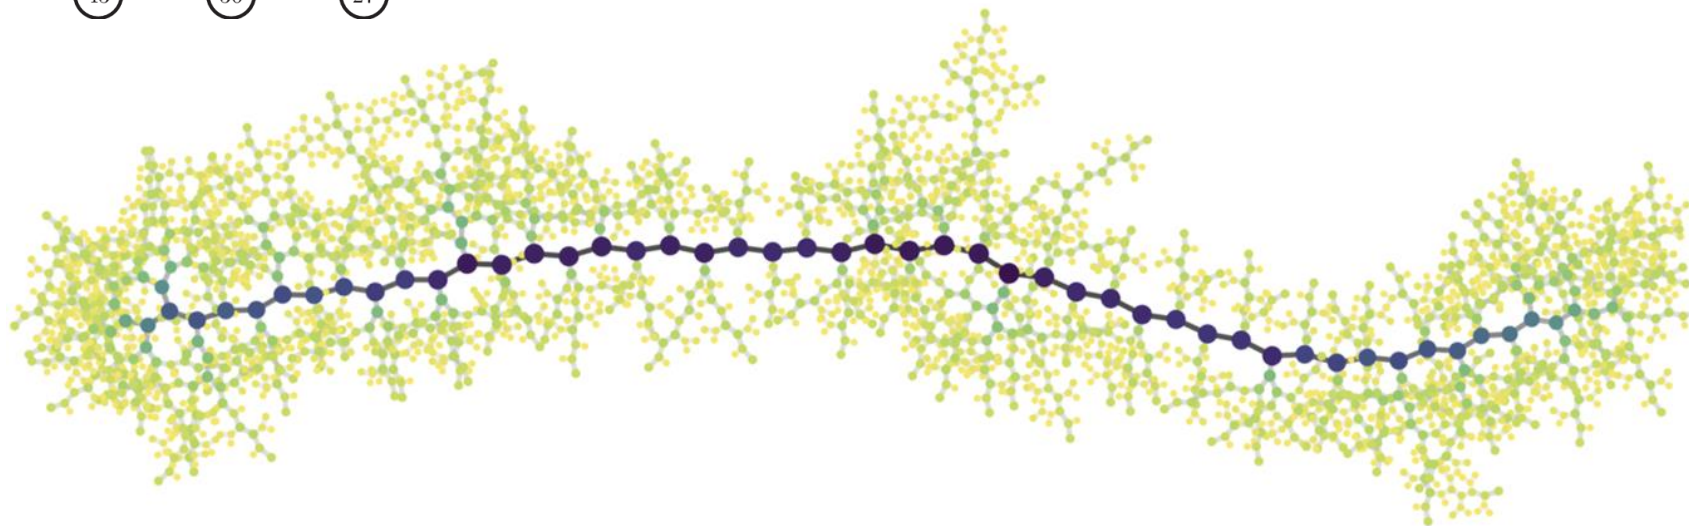
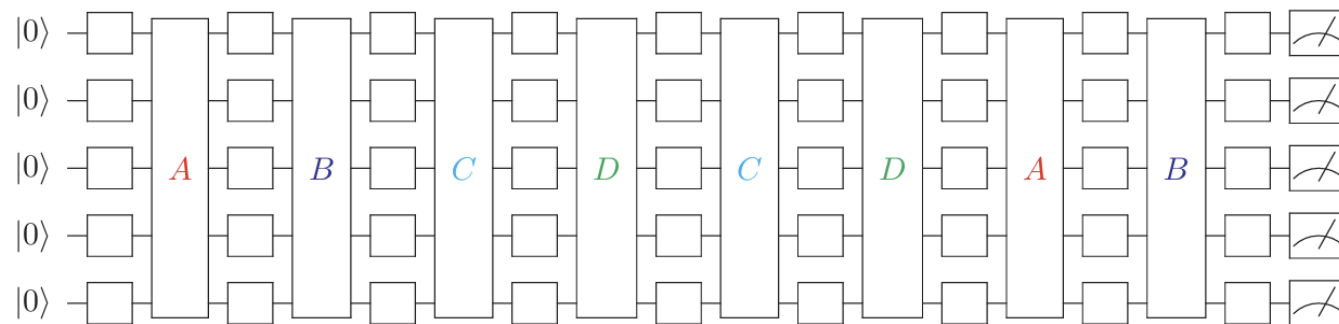
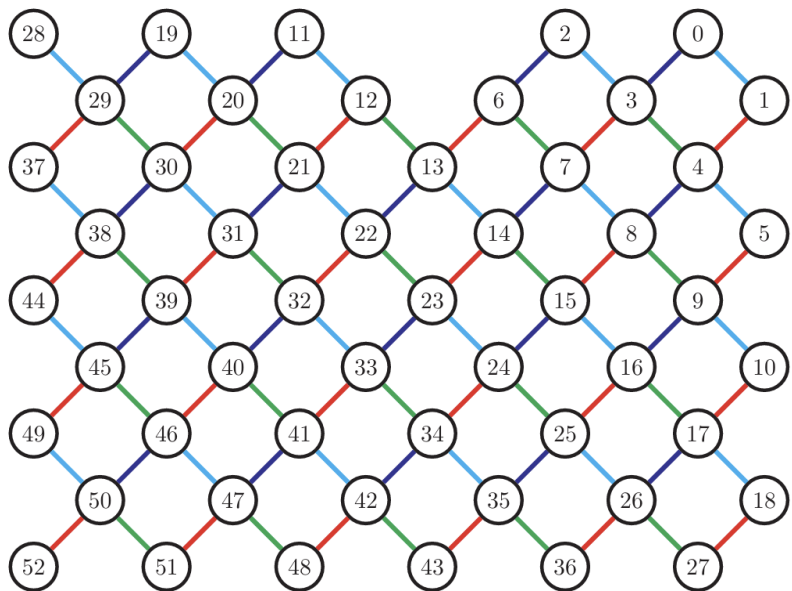
Application: Quantum circuit simulation



Theorem.

Let C be a quantum circuit with T gates and whose underlying circuit graph is G_C . Then C can be simulated deterministically in time $T^{O(1)} \exp(O(\text{tw}(G_C)))$.

Google's supremacy experiment



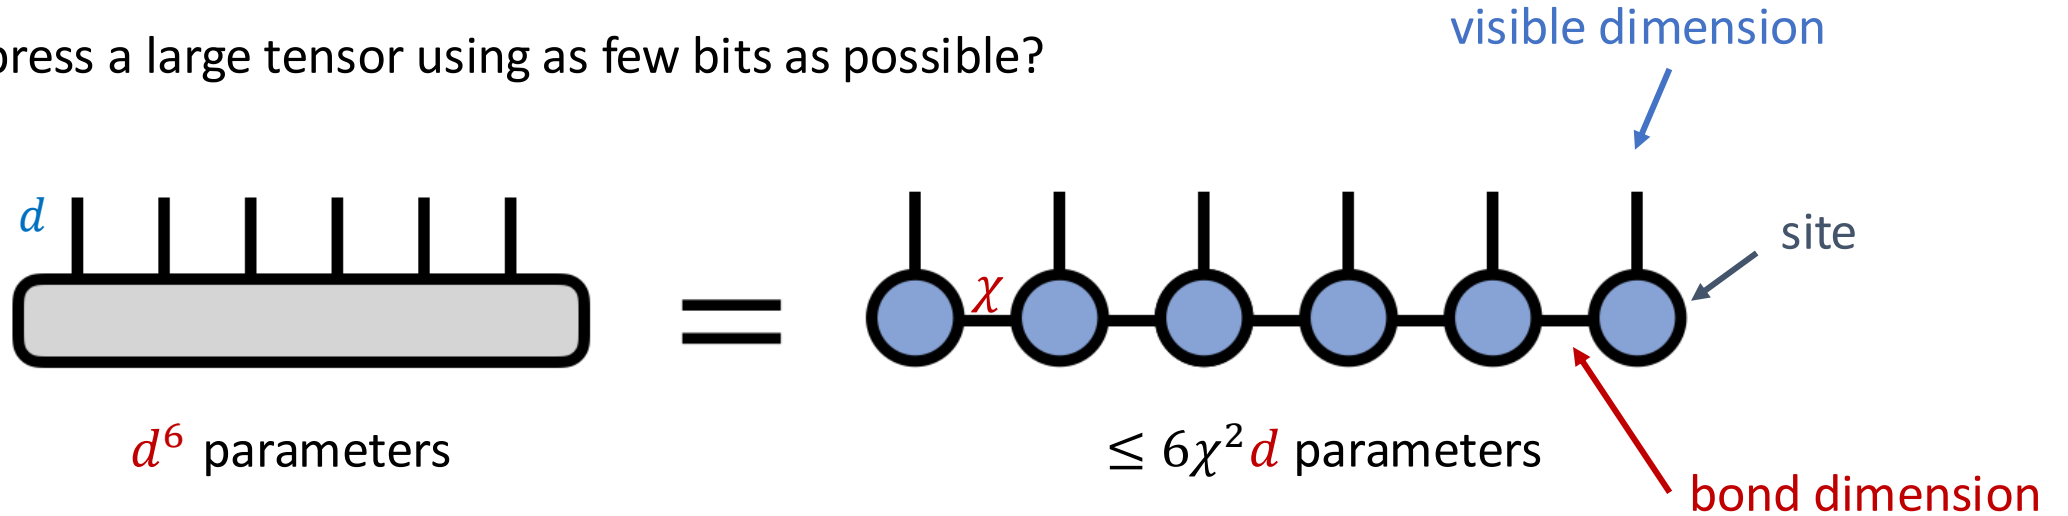
Tensor network algorithms



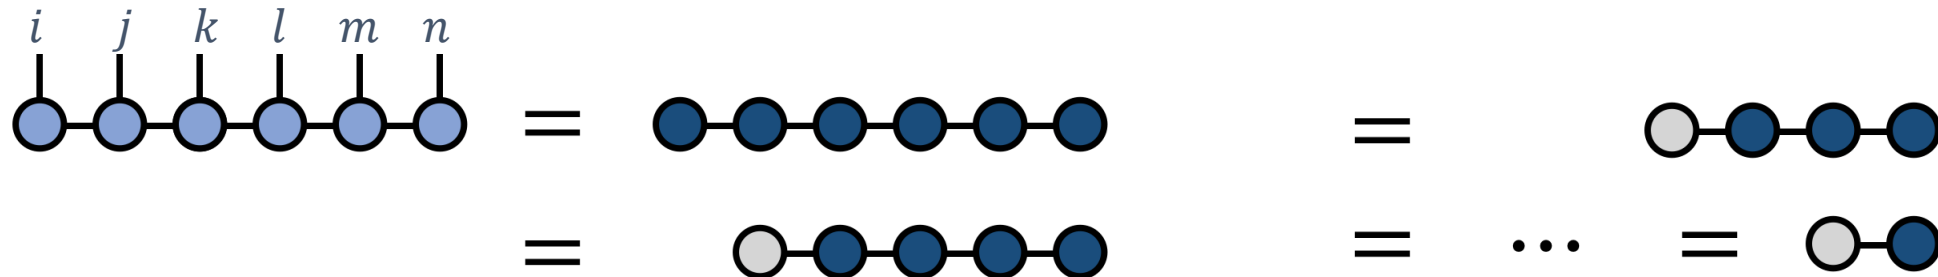
Matrix product states (MPS) / tensor trains



How to use express a large tensor using as few bits as possible?

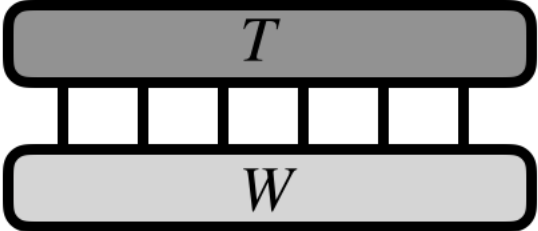


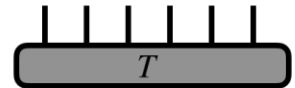
$$T_{ijklmn} = \sum_{\alpha_1, \dots, \alpha_5} A_{\alpha_1}^i A_{\alpha_1 \alpha_2}^j A_{\alpha_2 \alpha_3}^k A_{\alpha_3 \alpha_4}^l A_{\alpha_4 \alpha_5}^m A_{\alpha_5}^n$$

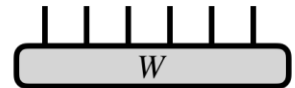


Cost: $\mathcal{O}(N\chi^2)$

Inner product of two MPSs

$$\langle T, W \rangle =$$


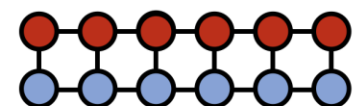
$$T = A^{s_1} A^{s_2} A^{s_3} A^{s_4} A^{s_5} A^{s_6}$$


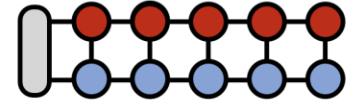
$$W = B^{s_1} B^{s_2} B^{s_3} B^{s_4} B^{s_5} B^{s_6}$$


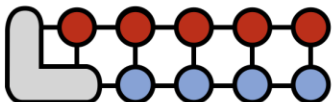
Contracting a
ladder TN

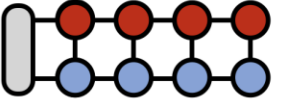


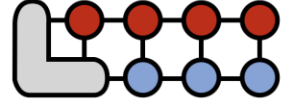
Cost: $\mathcal{O}(N\chi^3 d)$

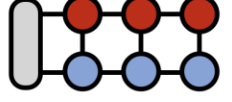



$$=$$


$$=$$


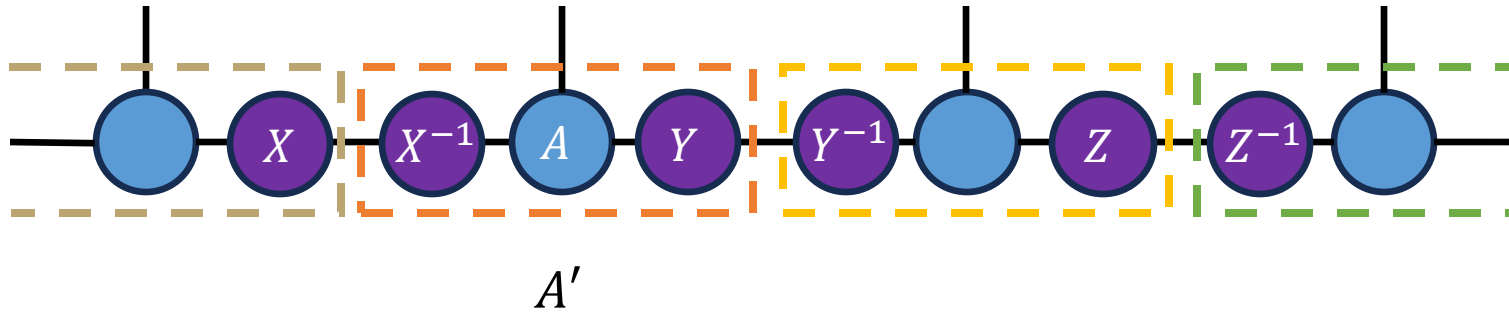
$$=$$


$$=$$


$$=$$


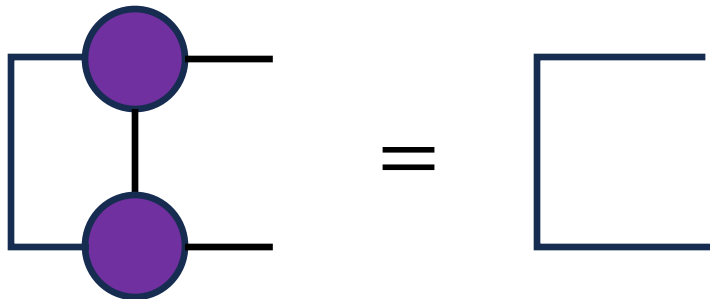
$$= \dots =$$


Symmetry in MPS



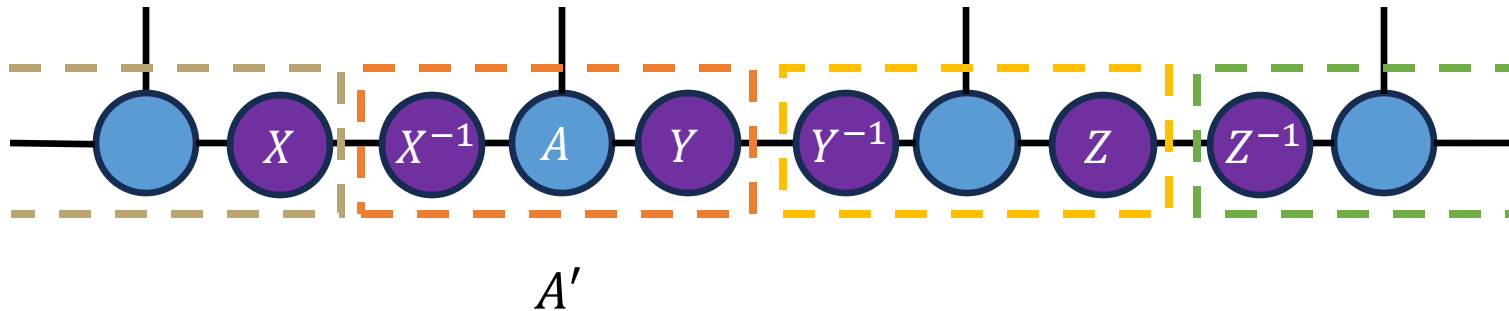
- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Left) canonical form:



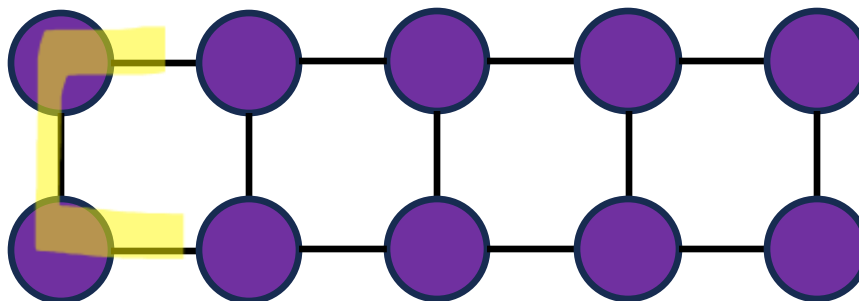
$$\sum_{i \in [d]} (A_n^i)^\dagger A_n^i = I \quad \forall n \in [N-1]$$

Symmetry in MPS

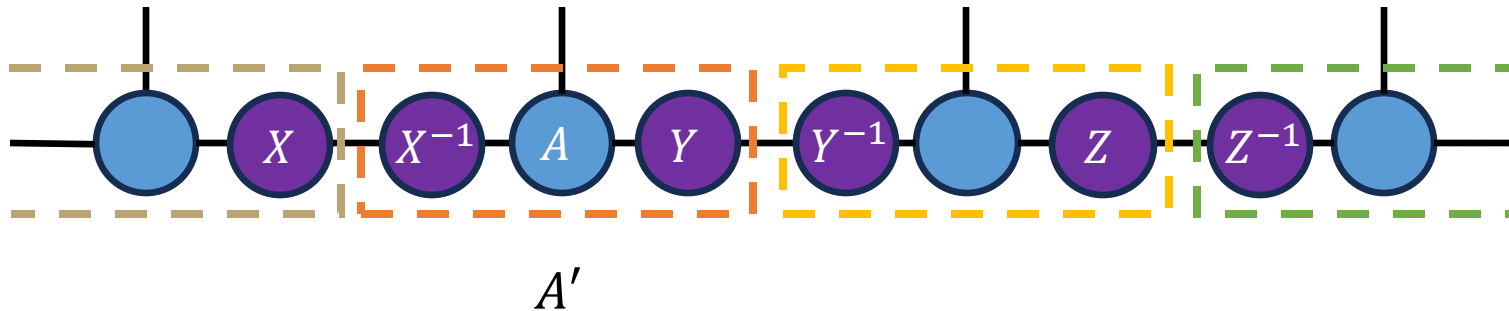


- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Left) canonical form:

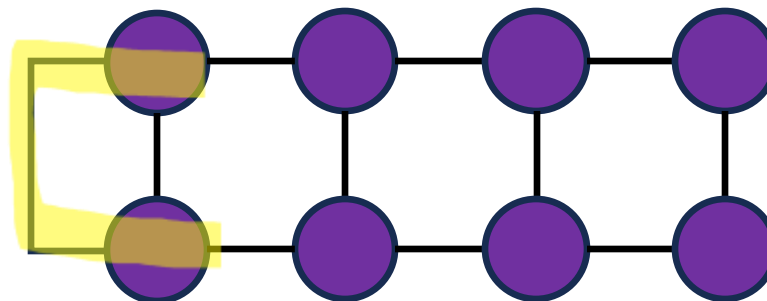


Symmetry in MPS

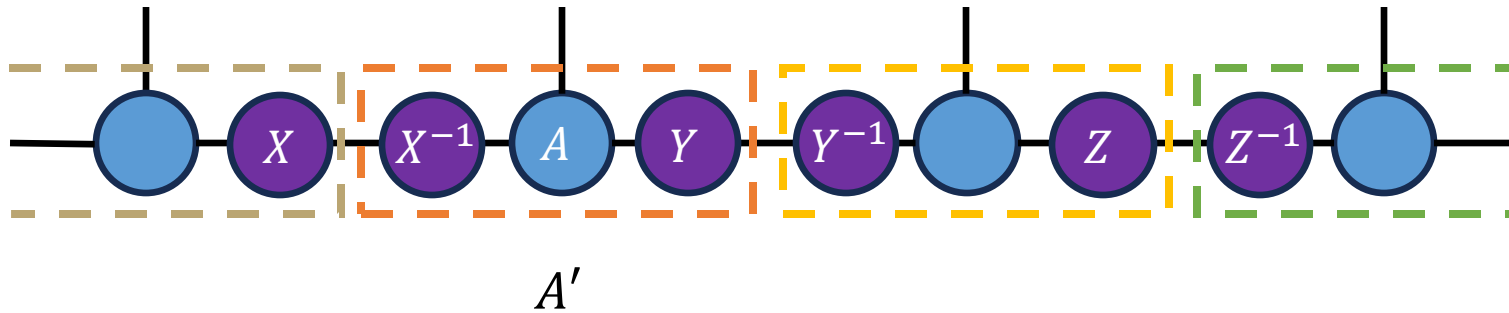


- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Left) canonical form:

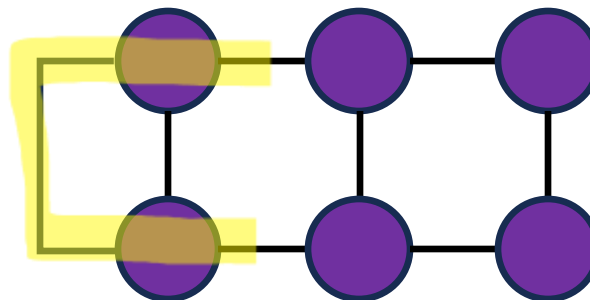


Symmetry in MPS

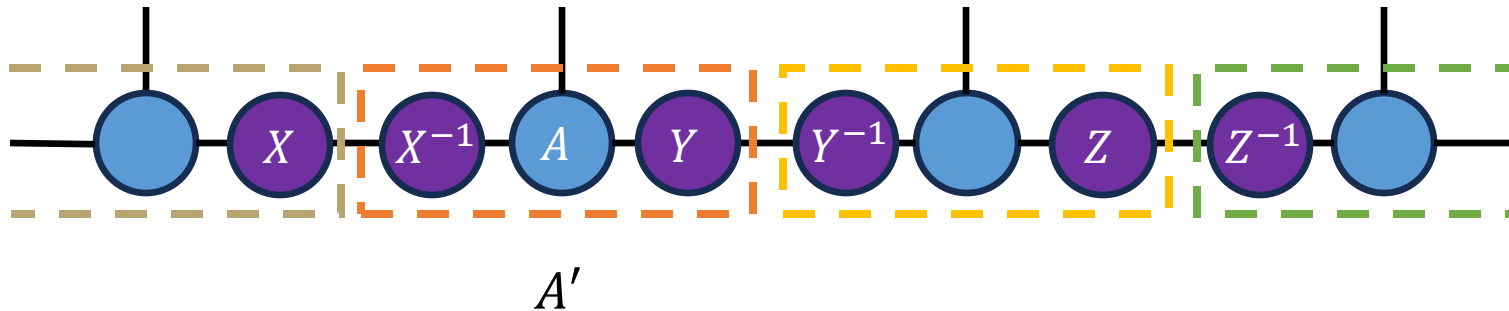


- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Left) canonical form:

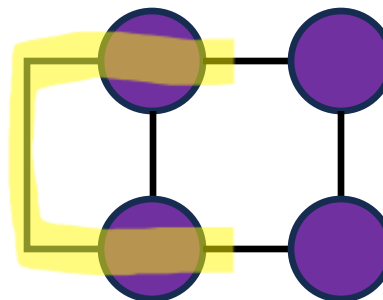


Symmetry in MPS

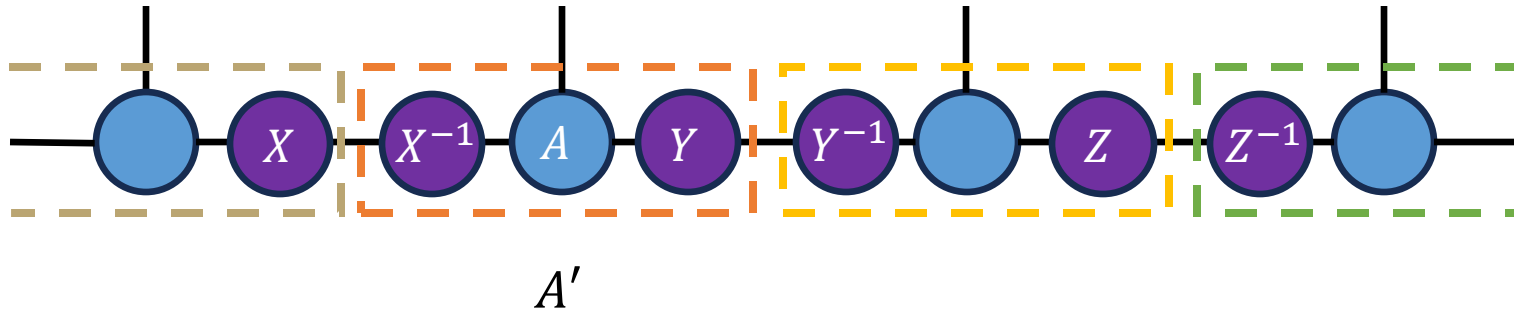


- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Left) canonical form:

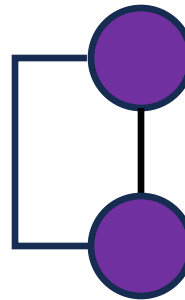


Symmetry in MPS



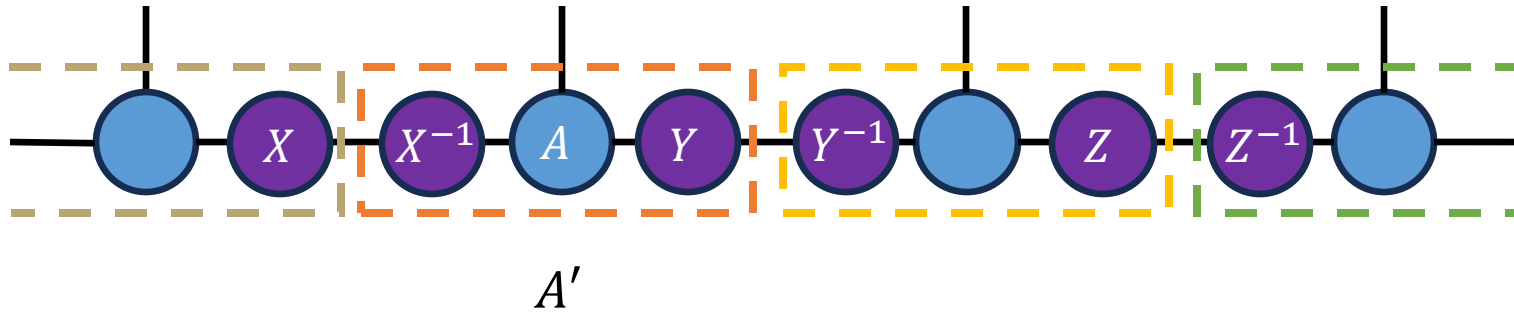
- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Left) canonical form:



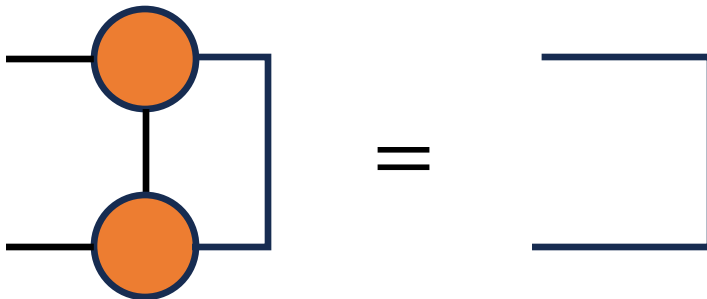
$$= \sum_{i \in [d]} \langle A_N^i, A_N^i \rangle = \|T\|_F^2$$

Symmetry in MPS



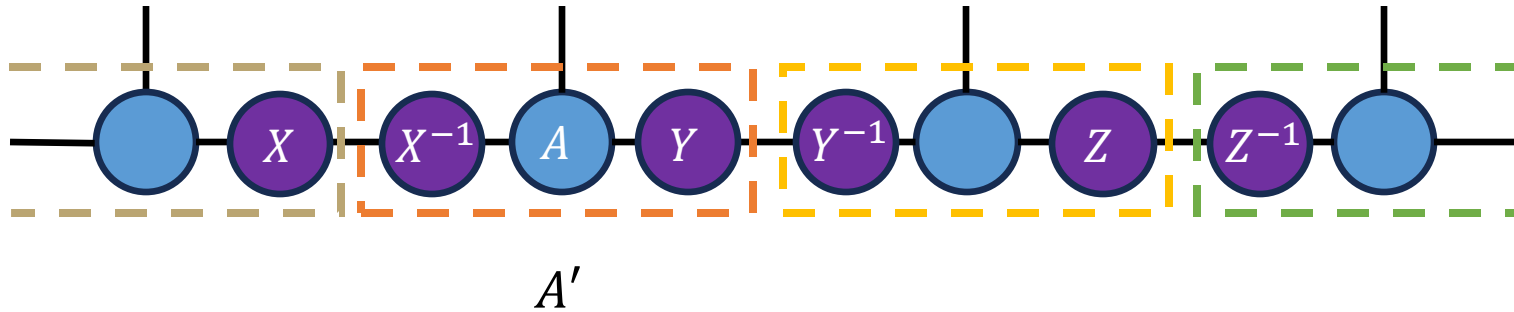
- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

(Right) canonical form:



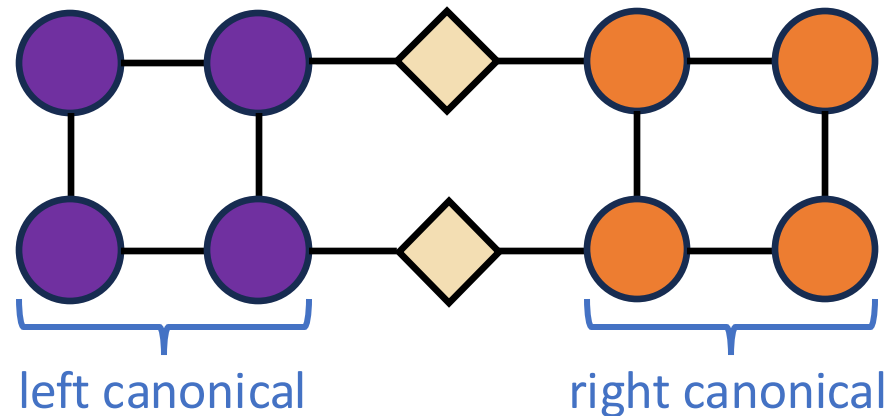
$$\sum_{i \in [d]} A_n^i (A_n^i)^\dagger = I \quad \forall n \in \{2, \dots, N\}$$

Symmetry in MPS



- Inserting invertible matrices X and X' does not change the whole tensor
- The MPS form is not unique

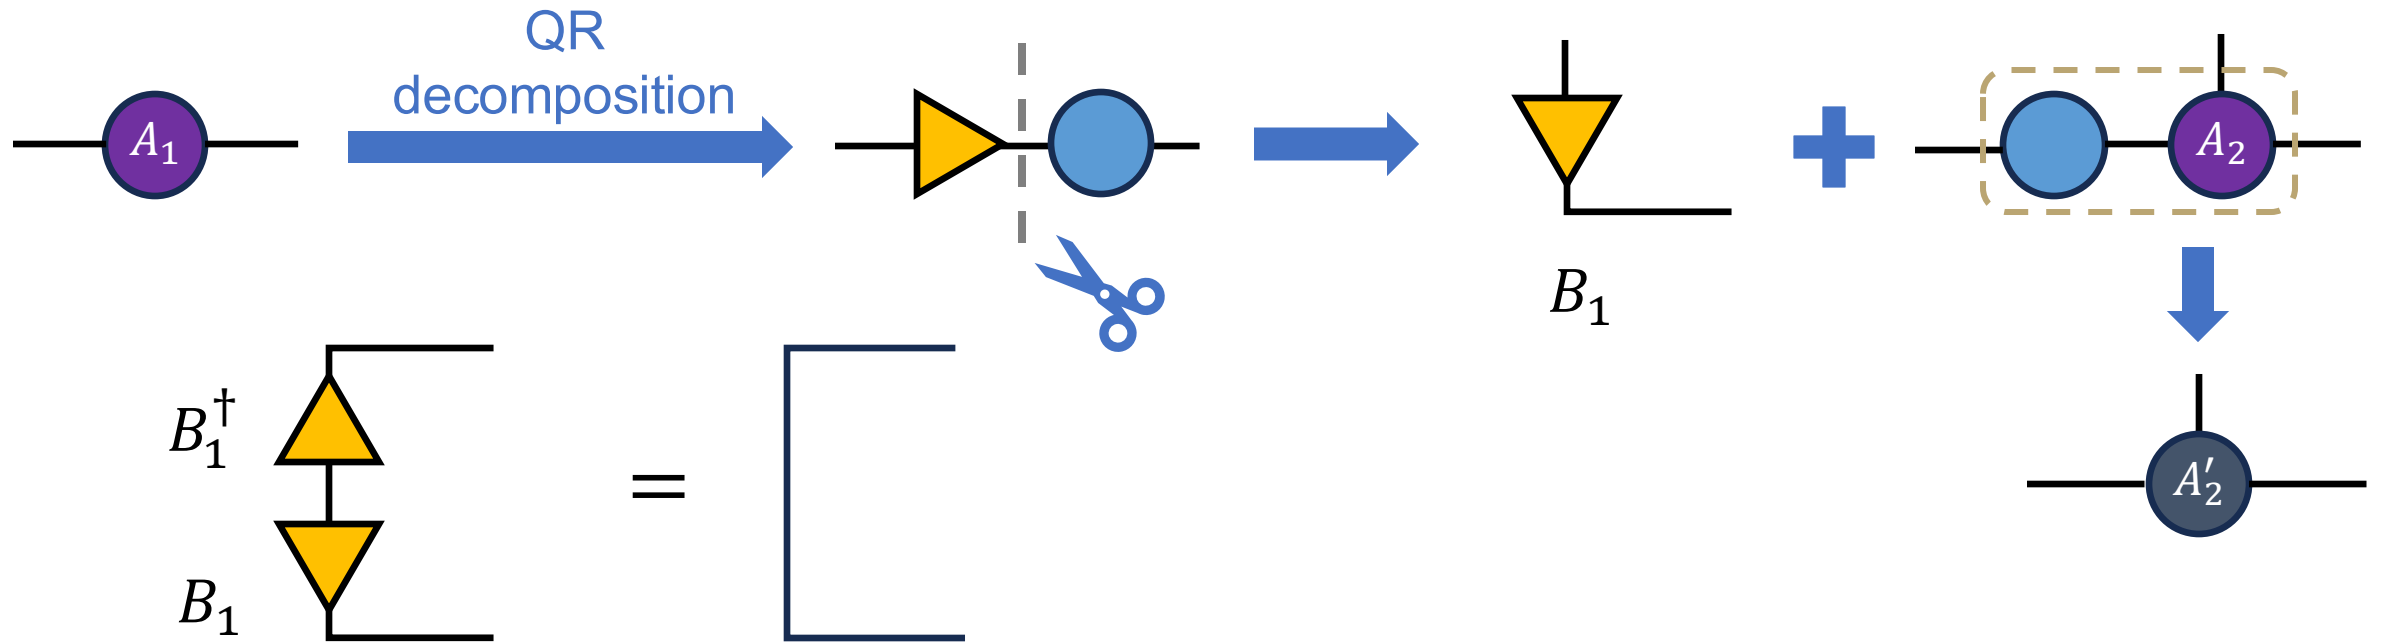
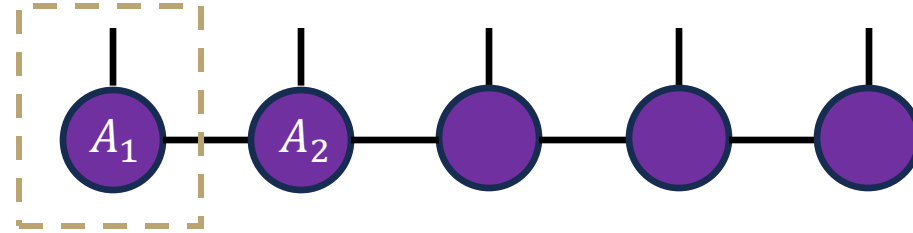
(Center) canonical form:



Canonicalize an MPS

Given an arbitrary MPS, how to make it in the (left/right/center) canonical form?

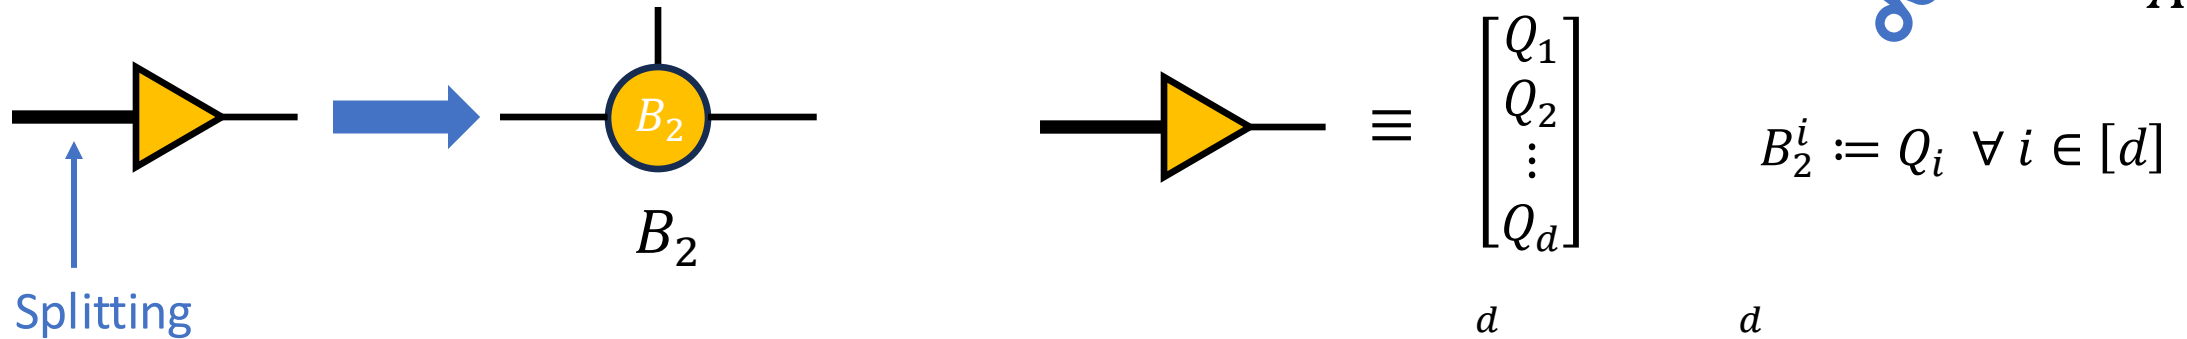
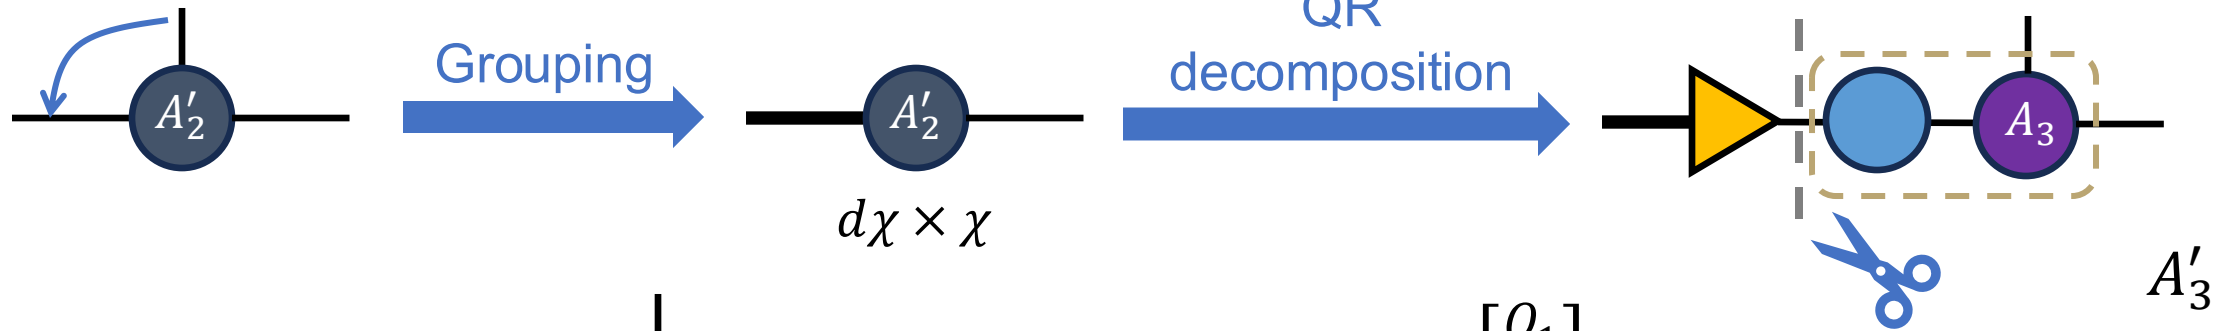
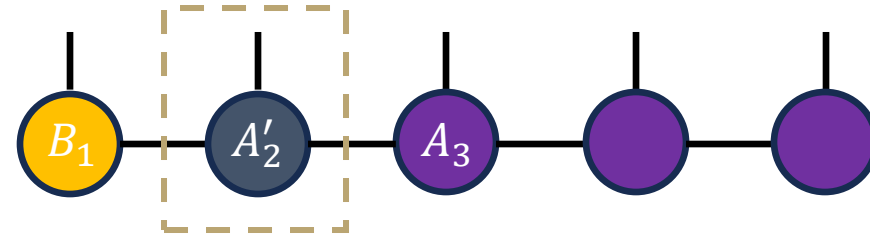
Short answer: QR!



Canonicalize an MPS

Given an arbitrary MPS, how to make it in the (left/right/center) canonical form?

Short answer: QR!

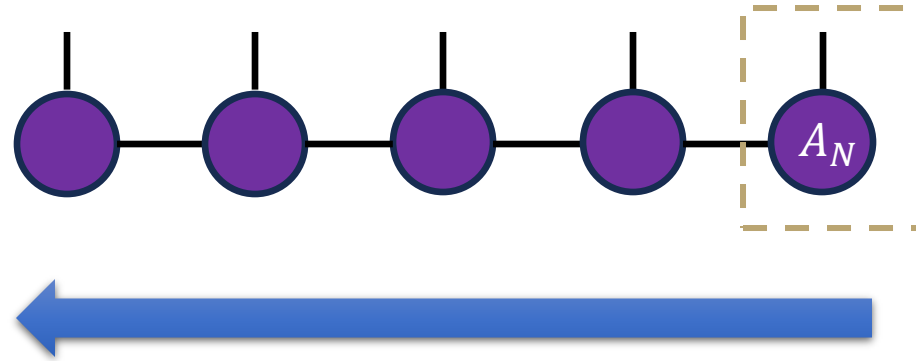


$$I = Q^\dagger Q = \sum_{i=1}^d Q_i^\dagger Q_i = \sum_{i=1}^d (B_2^i)^\dagger B_2^i$$

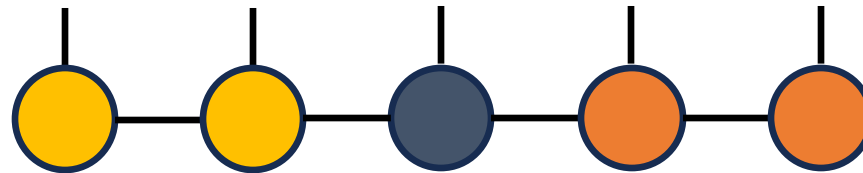
Canonicalize an MPS

Given an arbitrary MPS, how to make it in the (left/right/center) canonical form?

Short answer: QR!

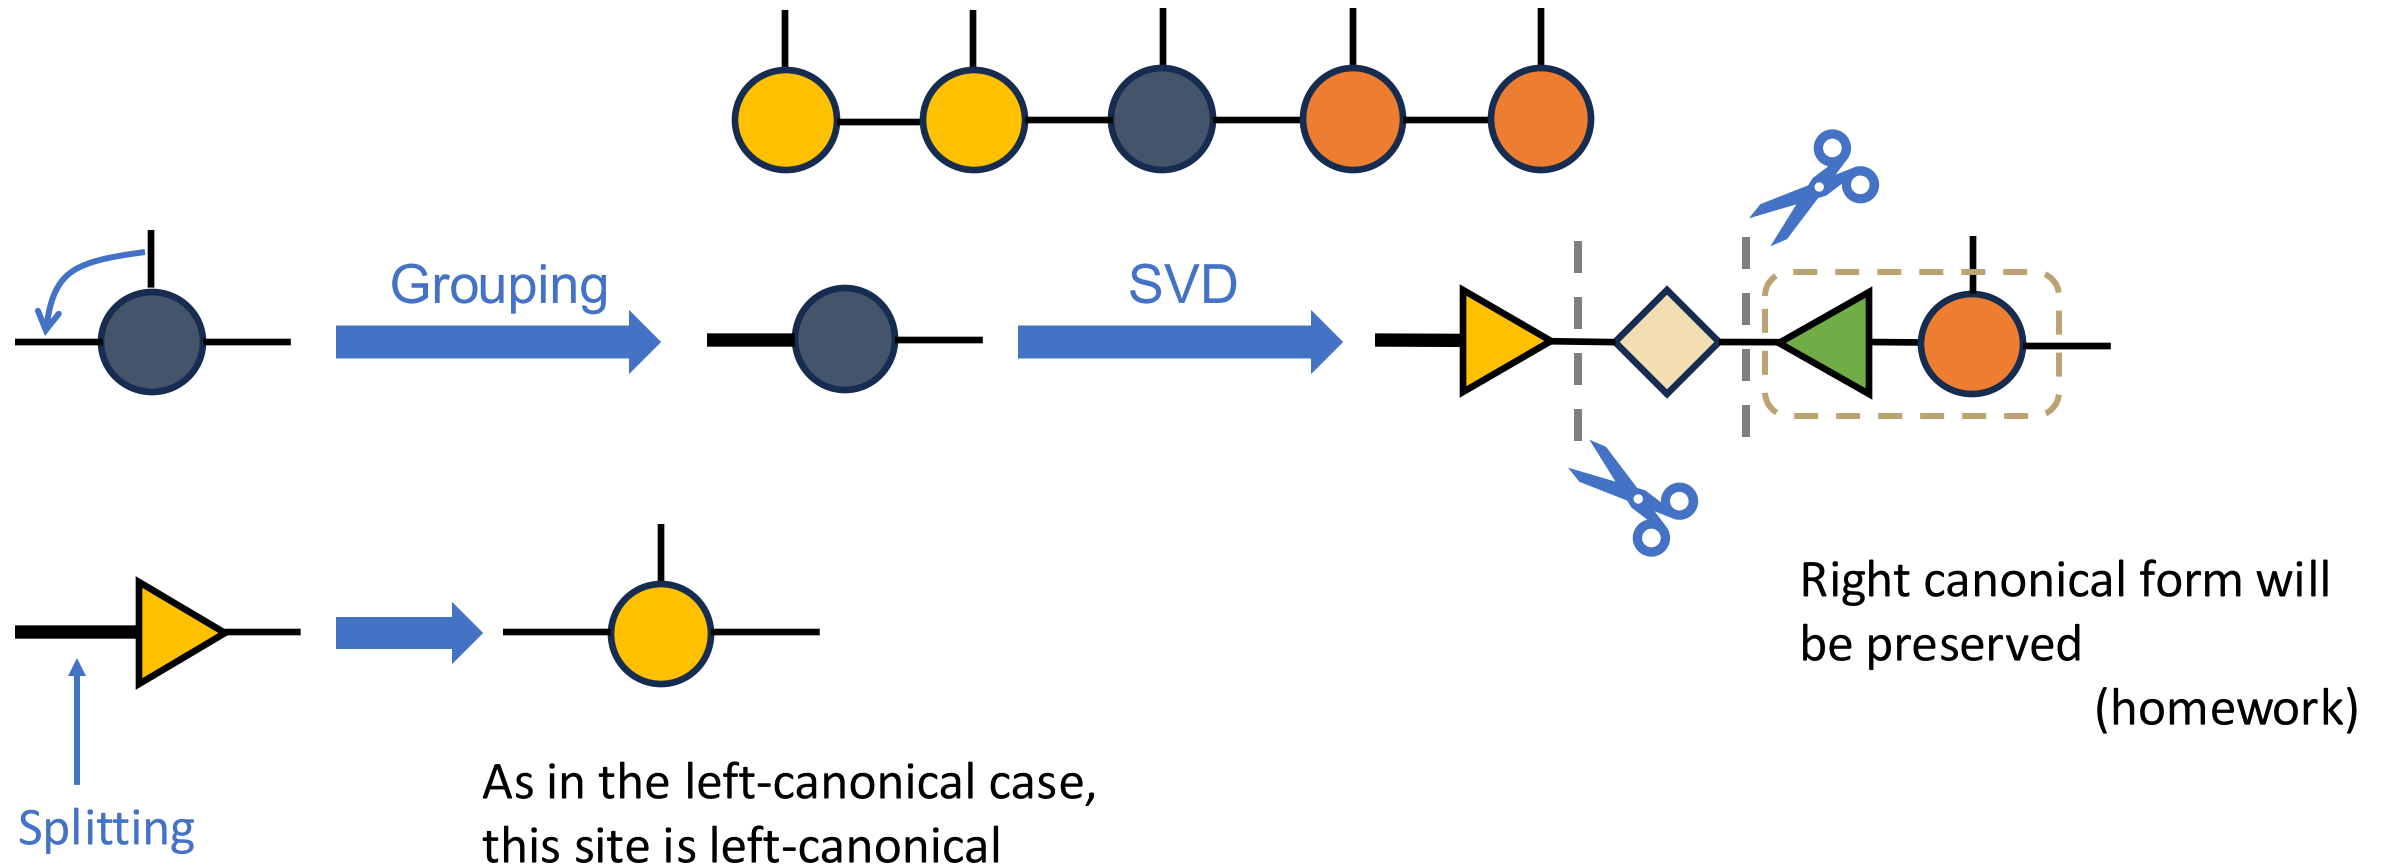


- Apply the same procedure from right to left to make the MPS **right canonical**
- For center canonical, we left-canonicalize the first half and right-canonicalize the second half



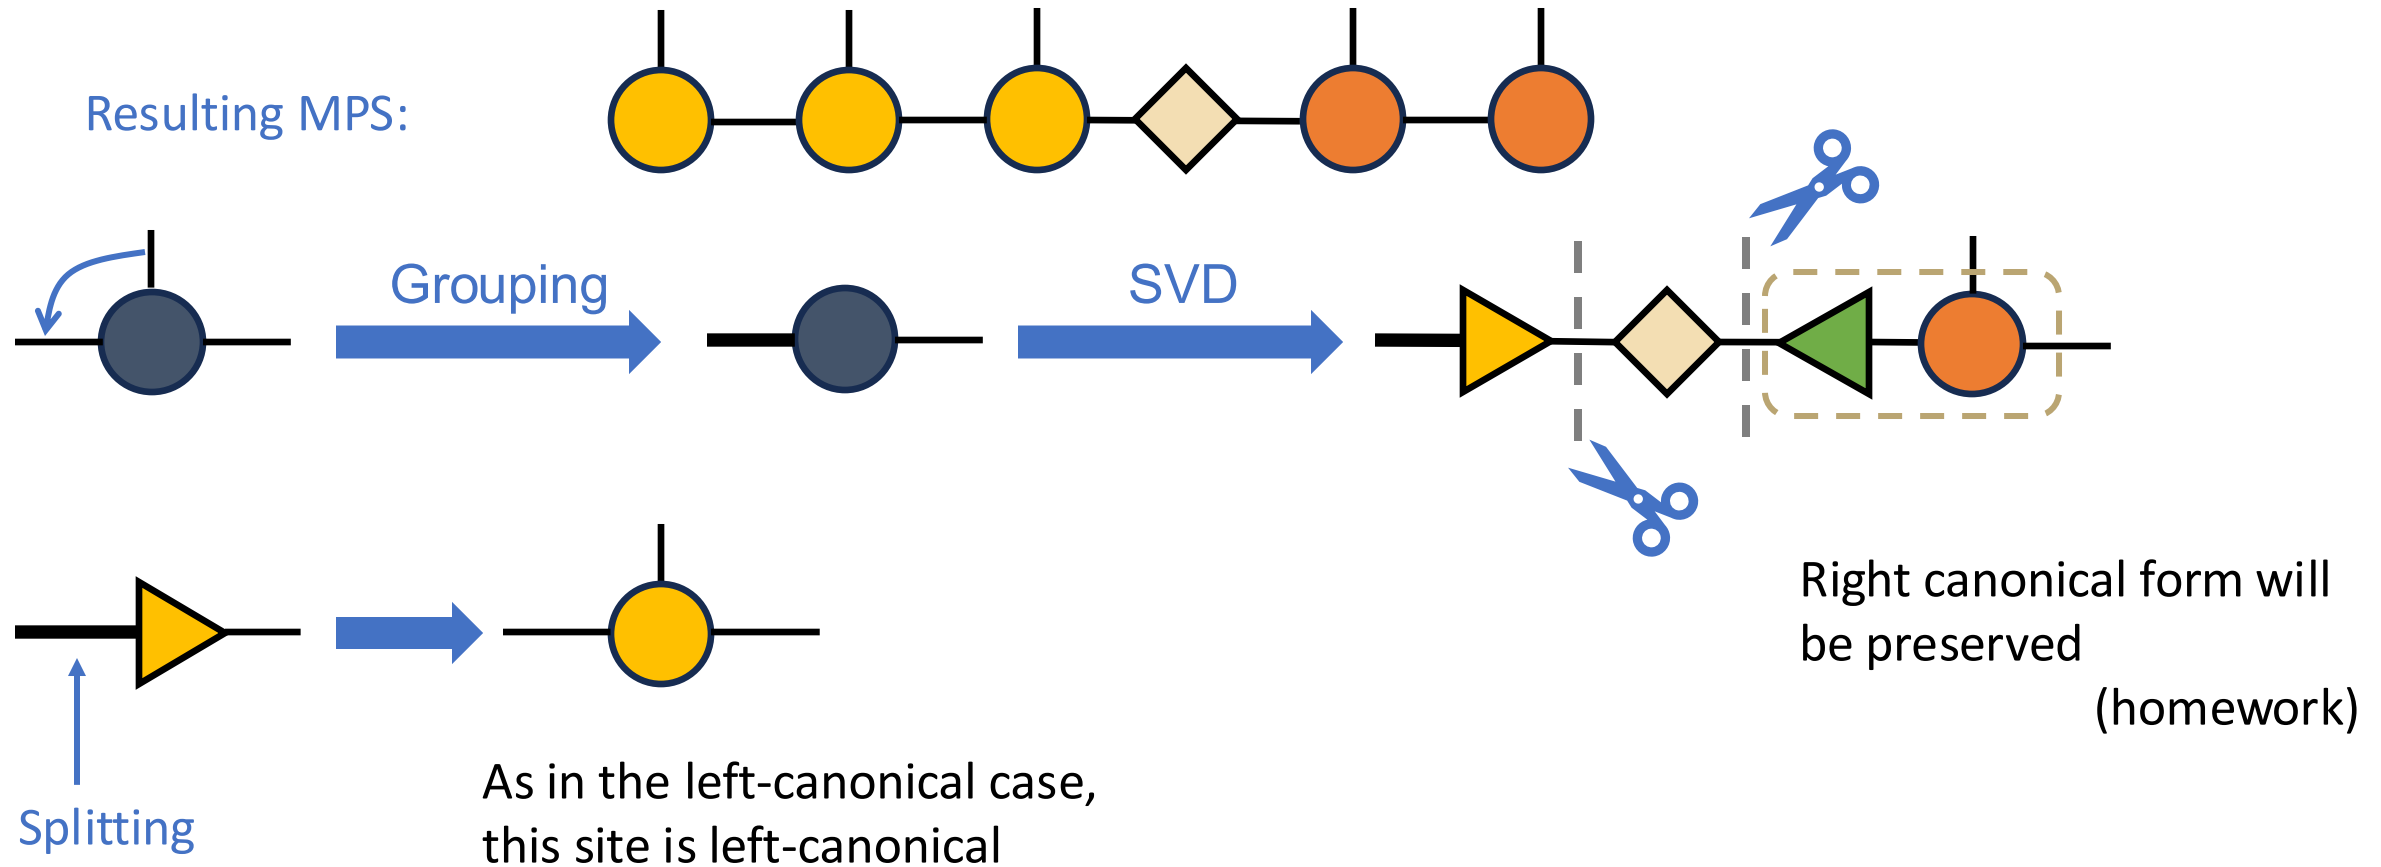
Center-canonicalize an MPS

For center canonical form, we left-canonicalize the first half and right-canonicalize the second half



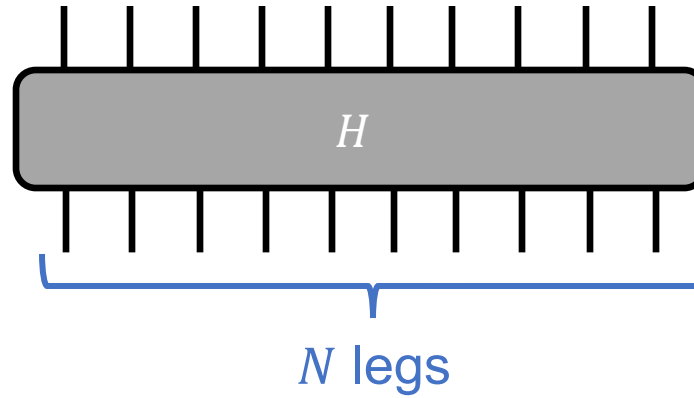
Center-canonicalize an MPS

For center canonical form, we left-canonicalize the first half and right-canonicalize the second half



DMRG: setup

We want to find the minimum eigenvalue of a **huge** matrix H (say 2^N -by- 2^N)

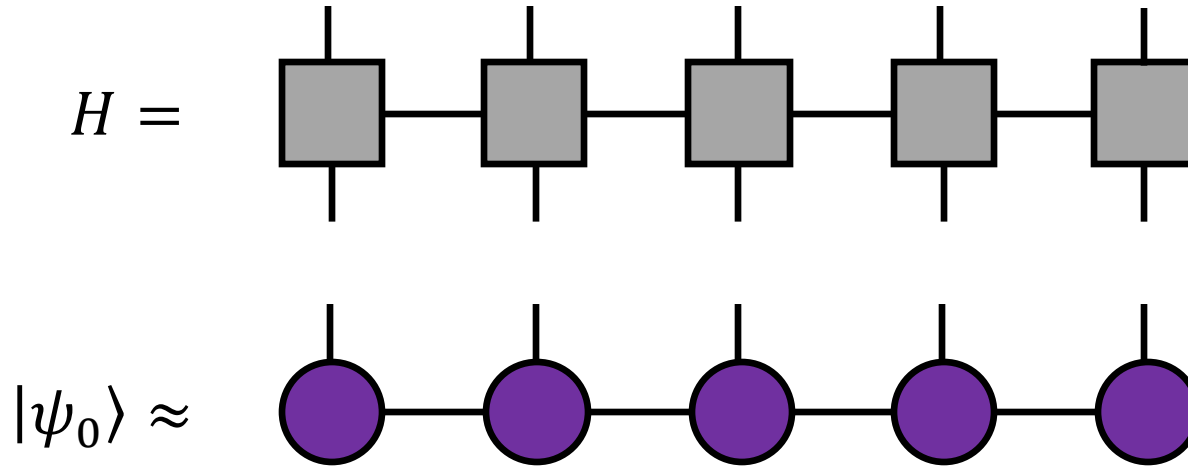


DMRG: setup

We want to find the **minimum eigenvalue** of a **huge** matrix H (say 2^N -by- 2^N)

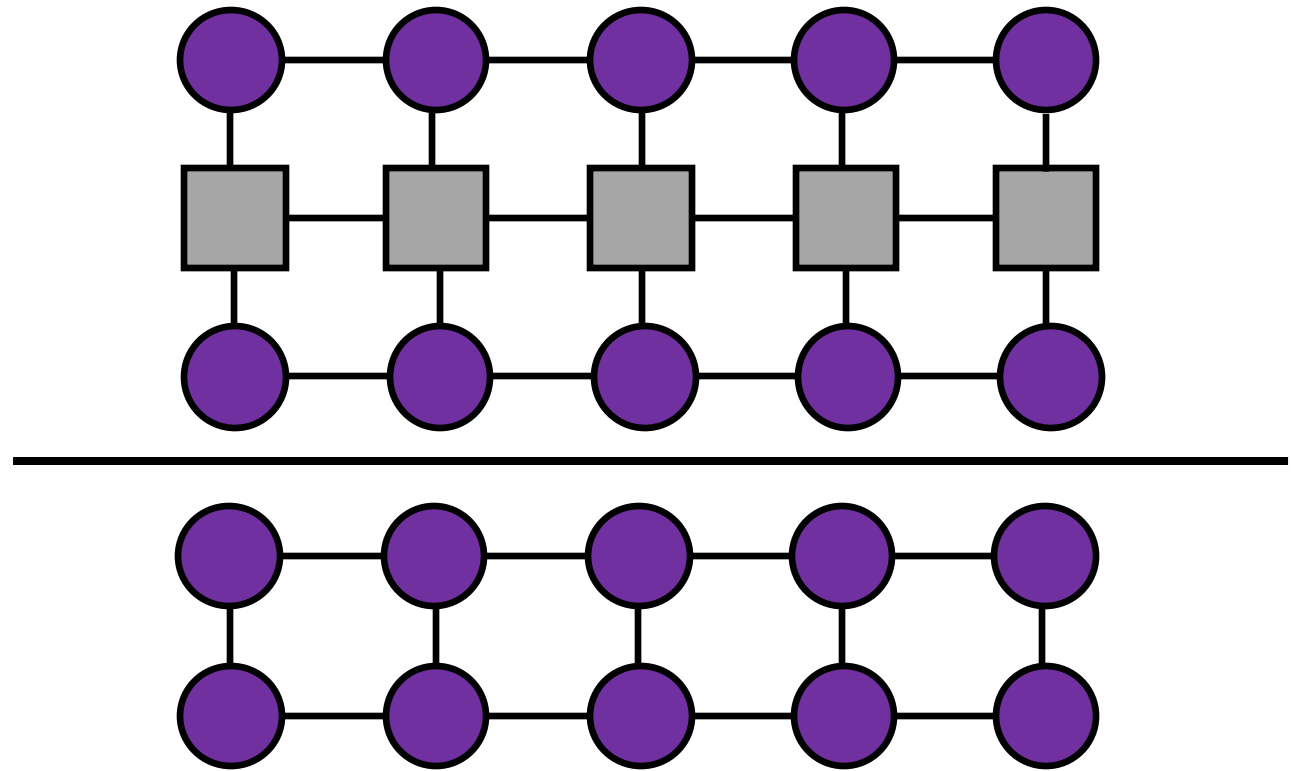
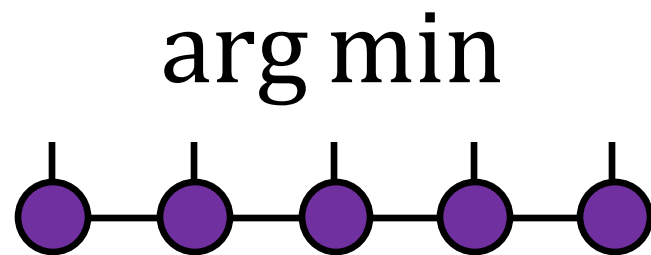
Assumptions:

- H has a “low-rank” representation; more formally, H is a matrix product operator (MPO)
- The minimum eigenvector can also be (approximately) represented as an MPS



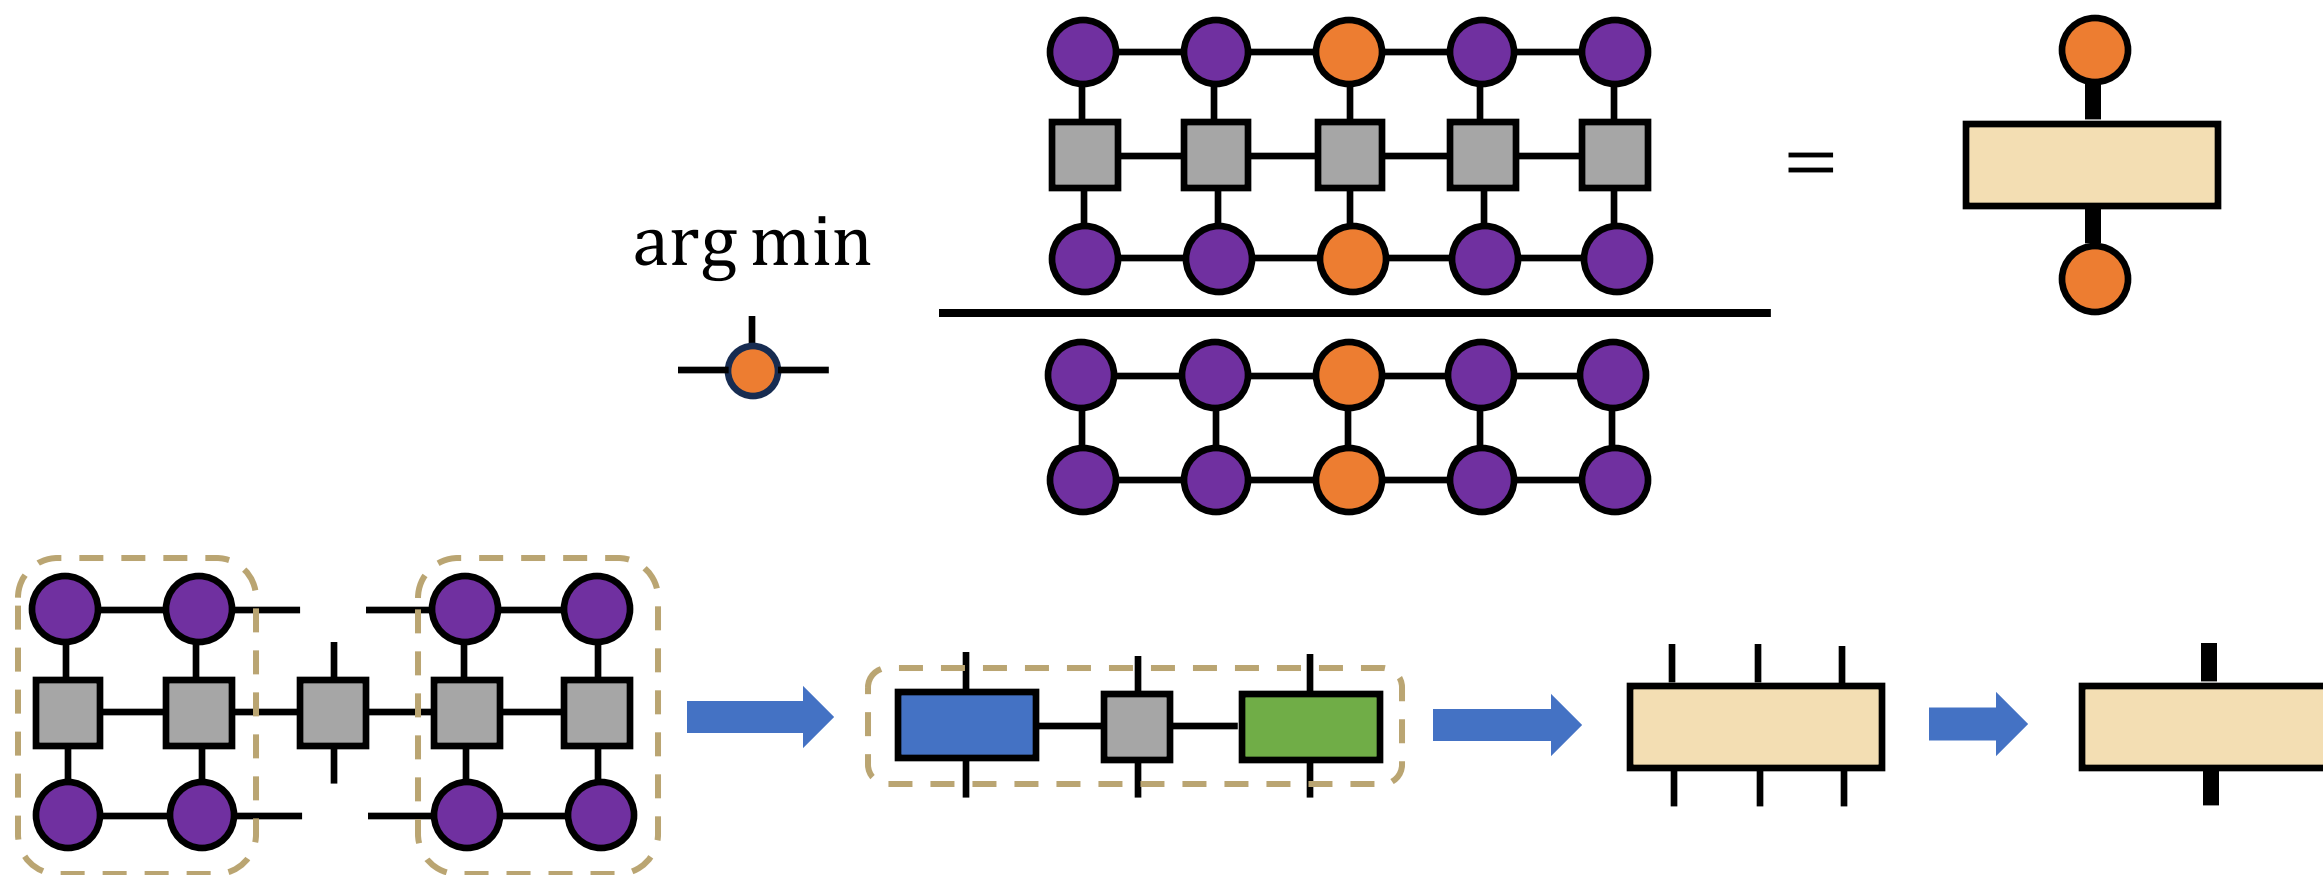
DMRG: setup

How to solve this non-linear optimization problem?



DMRG: algorithm

Core idea: **alternating minimization!**

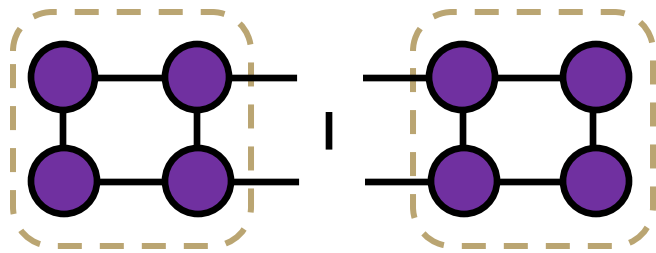
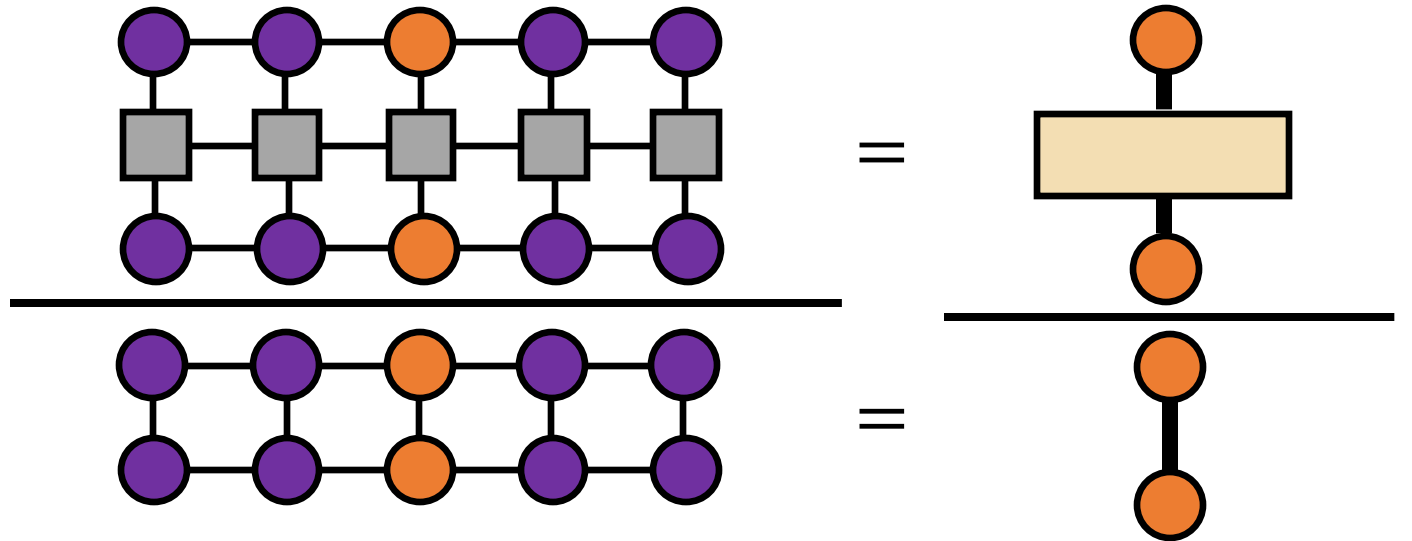


DMRG: algorithm

Core idea: **alternating minimization!**

➡ **Low-dimensional** ($2\chi^2$ -dim) eigenvalue problem

arg min



if center canonical



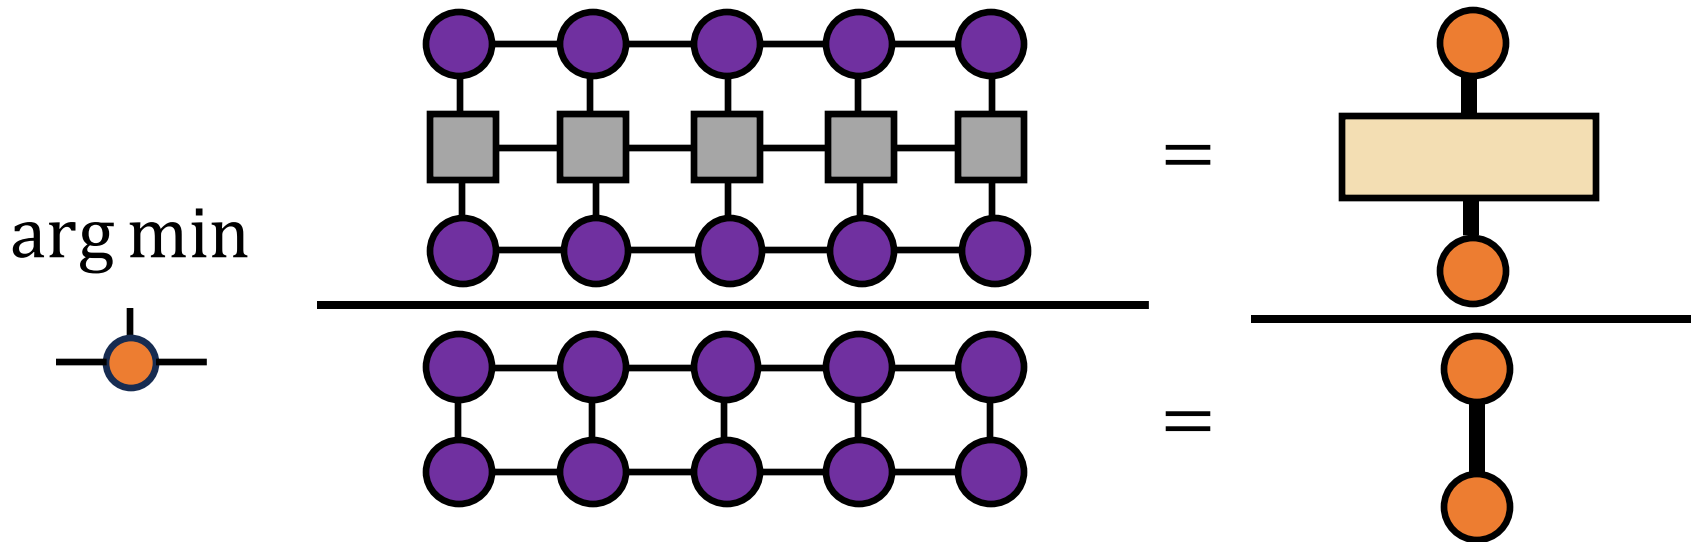
DMRG: algorithm

Core idea: **alternating minimization!**



Low-dimensional ($2\chi^2$ -dim) eigenvalue problem

left sweep + right sweep + left sweep + right sweep + \dots (until converge)



DMRG: applications

- This is a purely **heuristic** approach, but **super powerful** in practice!

Density matrix formulation for quantum renormalization groups



SR White - Physical review letters, 1992



A generalization of the numerical renormalization-group procedure used first by Wilson for the Kondo problem is presented. It is shown that this formulation is optimal in a certain sense. As a demonstration of the effectiveness of this approach, results from numerical real-space renormalization-group calculations for Heisenberg chains are presented.

☆  Cite **Cited by 9582** Related articles All 8 versions

[PDF] aps.org

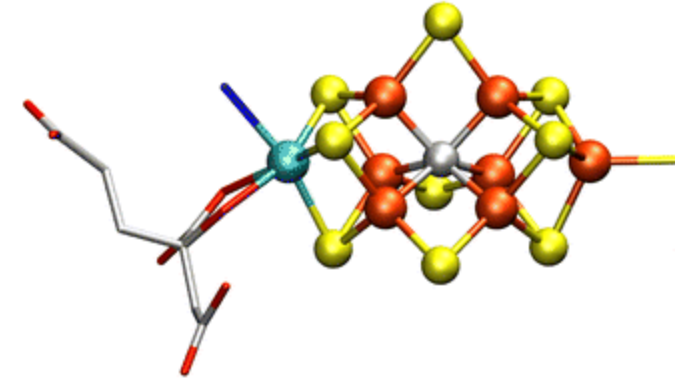


 Scholar **About 27,400 results** (0.11 sec) YEAR 

DMRG: applications

- This is a purely **heuristic** approach, but **super powerful** in practice!
- Widely used in quantum many-body physics and quantum chemistry
- Quantum language:
 - H : the **Hamiltonian**, which describes the dynamics of a physical system
 - The minimum eigenvalue: **ground state energy**
 - The minimum eigenvector: ground state
 - These capture key properties of the system at very low temperature
- **When DMRG is effective?**
 - H is an MPO
 - The ground state can be (approximately) represented as an MPS with low bond dimension



} 1D quantum many-body system

DMRG for 1D system

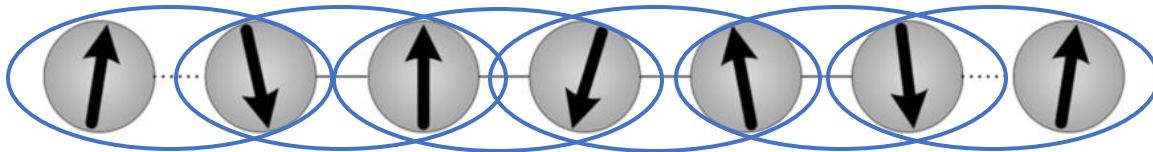
- Classical 1D spin chain (Ising model):



$$H_{\text{Ising}}(\sigma) = - \sum_i J_i \sigma_i \sigma_{i+1}$$

$$\sigma \in \{-1, 1\}^n$$

- Quantum 1D system:



$$H = \sum_i h_{i,i+1}$$

$h_{i,i+1}$ only acts on qubits i and $i + 1$

$$\underbrace{I \otimes I \otimes \dots \otimes I}_{i-1} \otimes h \otimes \underbrace{I \otimes \dots \otimes I}_{n-i-1}$$

DMRG for 1D system

Theorem (Hastings, 2007; Arad-Kitaev-Landau-Vazirani, 2013).

Consider a 1D quantum spin chain of N sites, each of local dimension d , with a Hamiltonian $H = \sum_{i=1}^{N-1} h_{i,i+1}$ where each $h_{i,i+1}$ acts on nearest neighbors, $\|h_{i,i+1}\| \leq 1$, and the Hamiltonian has a **spectral gap** $\Delta > 0$ above the ground state.

Then, there exists an **MPS approximation** $|\psi_{\text{MPS}}\rangle$ to the ground state $|\psi_0\rangle$ such that

$$\| |\psi_{\text{MPS}}\rangle - |\psi_0\rangle \| \leq 1/\text{poly}(n)$$

with bond dimension $\chi = \exp\left(\mathcal{O}\left(\Delta^{-\frac{1}{3}} \log^{\frac{2}{3}} n\right)\right)$

- **Landau-Vazirani-Vidick; Arad-Landau-Vazirani-Vidick:** polynomial-time algorithms